BDM1213 Data Encoding Principles

Week 03: Big data and Apache Hadoop

Dr. James Hong

Scenario:
Bob has opened a small restaurant in his city

Traditional Scenario:

2 orders per hour

Single Cook

Food Shelf

Traditional Scenario:

Data is generated at a steady rate and is structured in nature

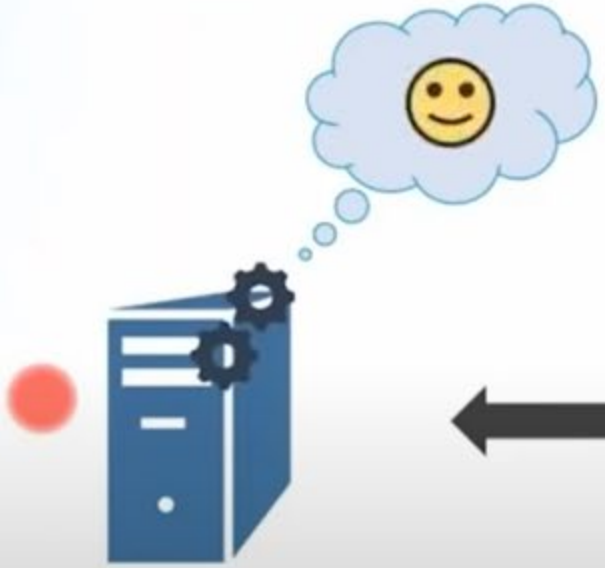

Traditional Processing System

RDBMS

Scenario 2:

- They started taking Online orders
- 10 orders per hour

Single Cook
(Regular Computing System)

Food Shelf
(Data)

Big Data Scenario:

Heterogenous data is being generated at an alarming rate by multiple sources

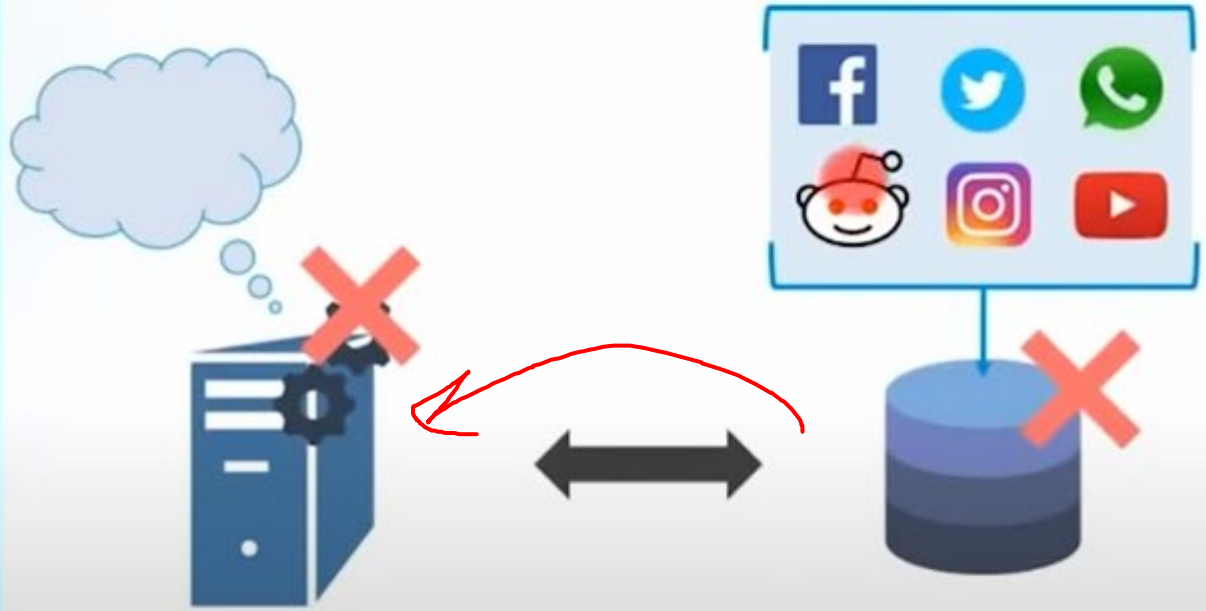

Traditional Processing System

RDBMS

Issue 1: Too Many Orders Per Hour

Solution: Hiring Multiple Cook

Scenario:

Multiple Cook cooking food

Issue:

Food Shelf becomes the BOTTLENECK

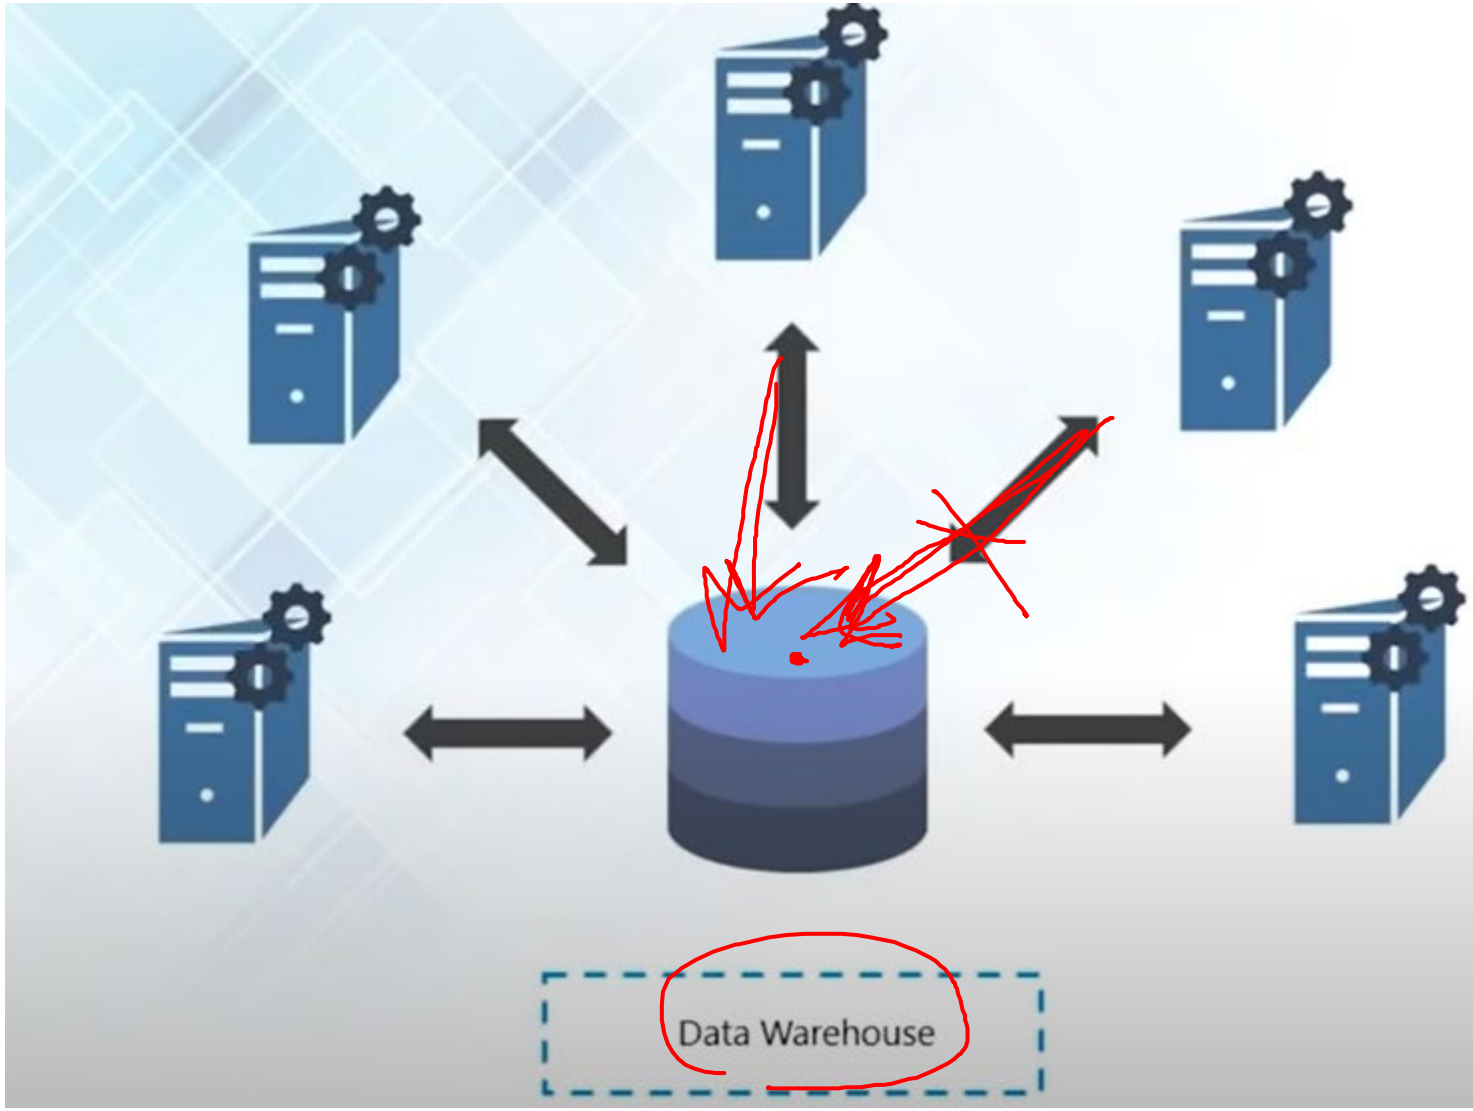

Scenario:

Multiple Processing Unit for data processing

Issue:

Bringing data to processing generated lots of Network overhead

Issue 2: Food Shelf becomes the Bottleneck

Solution: Distributed and Parallel Approach

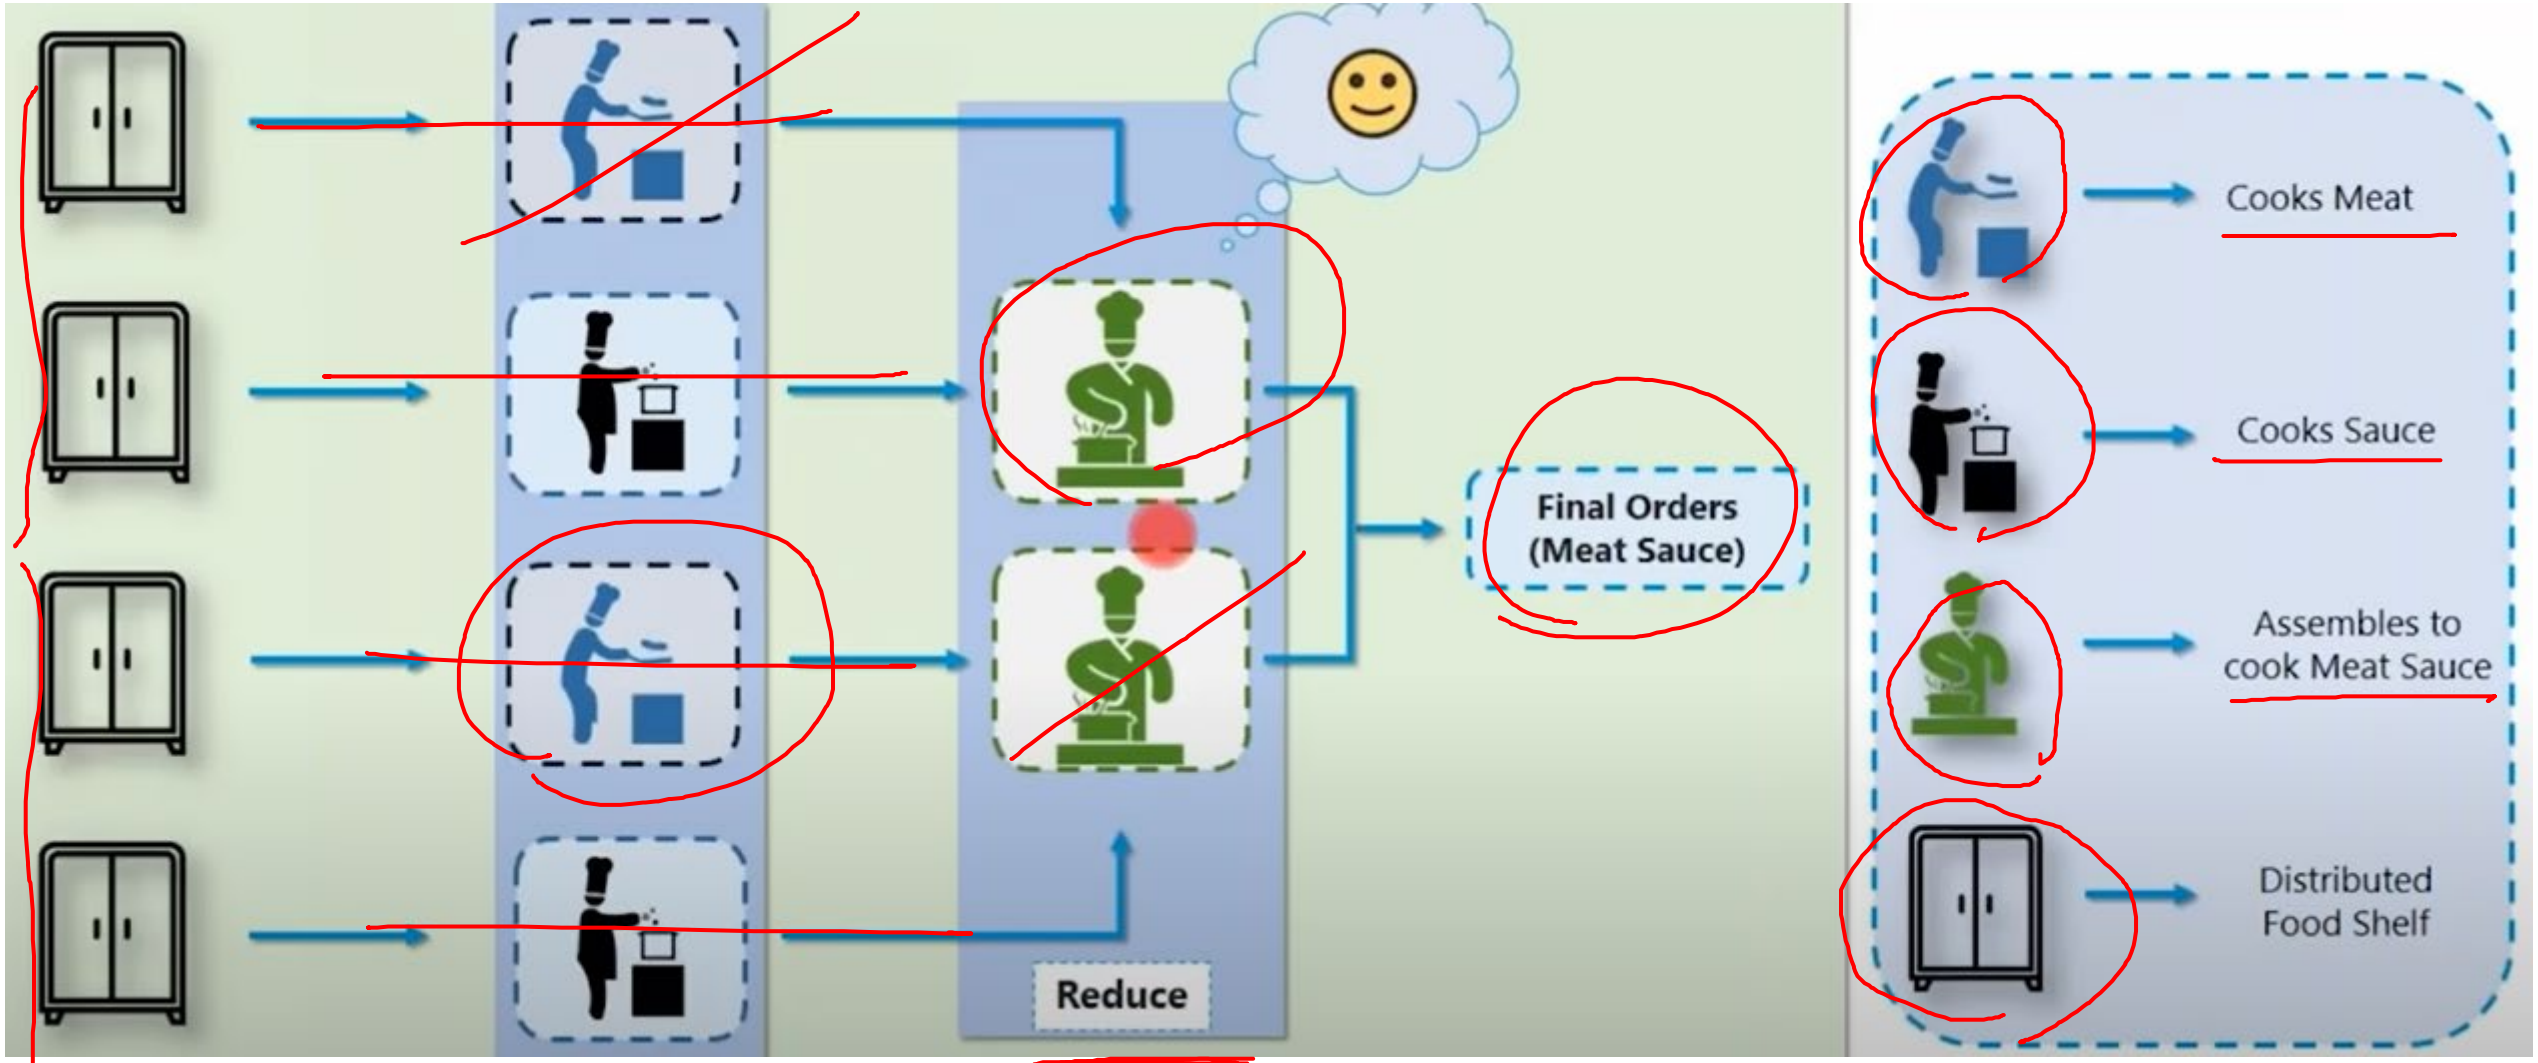
Final Orders
(Meat Sauce)
Reduce
Cooks Meat
Cooks Sauce
Assembles to
cook Meat Sauce
Distributed
Food Shelf

Map

Apache Hadoop:
Framework to Process Big Data

Hadoop is a framework that allows us to store and process large data sets in parallel and distributed fashion

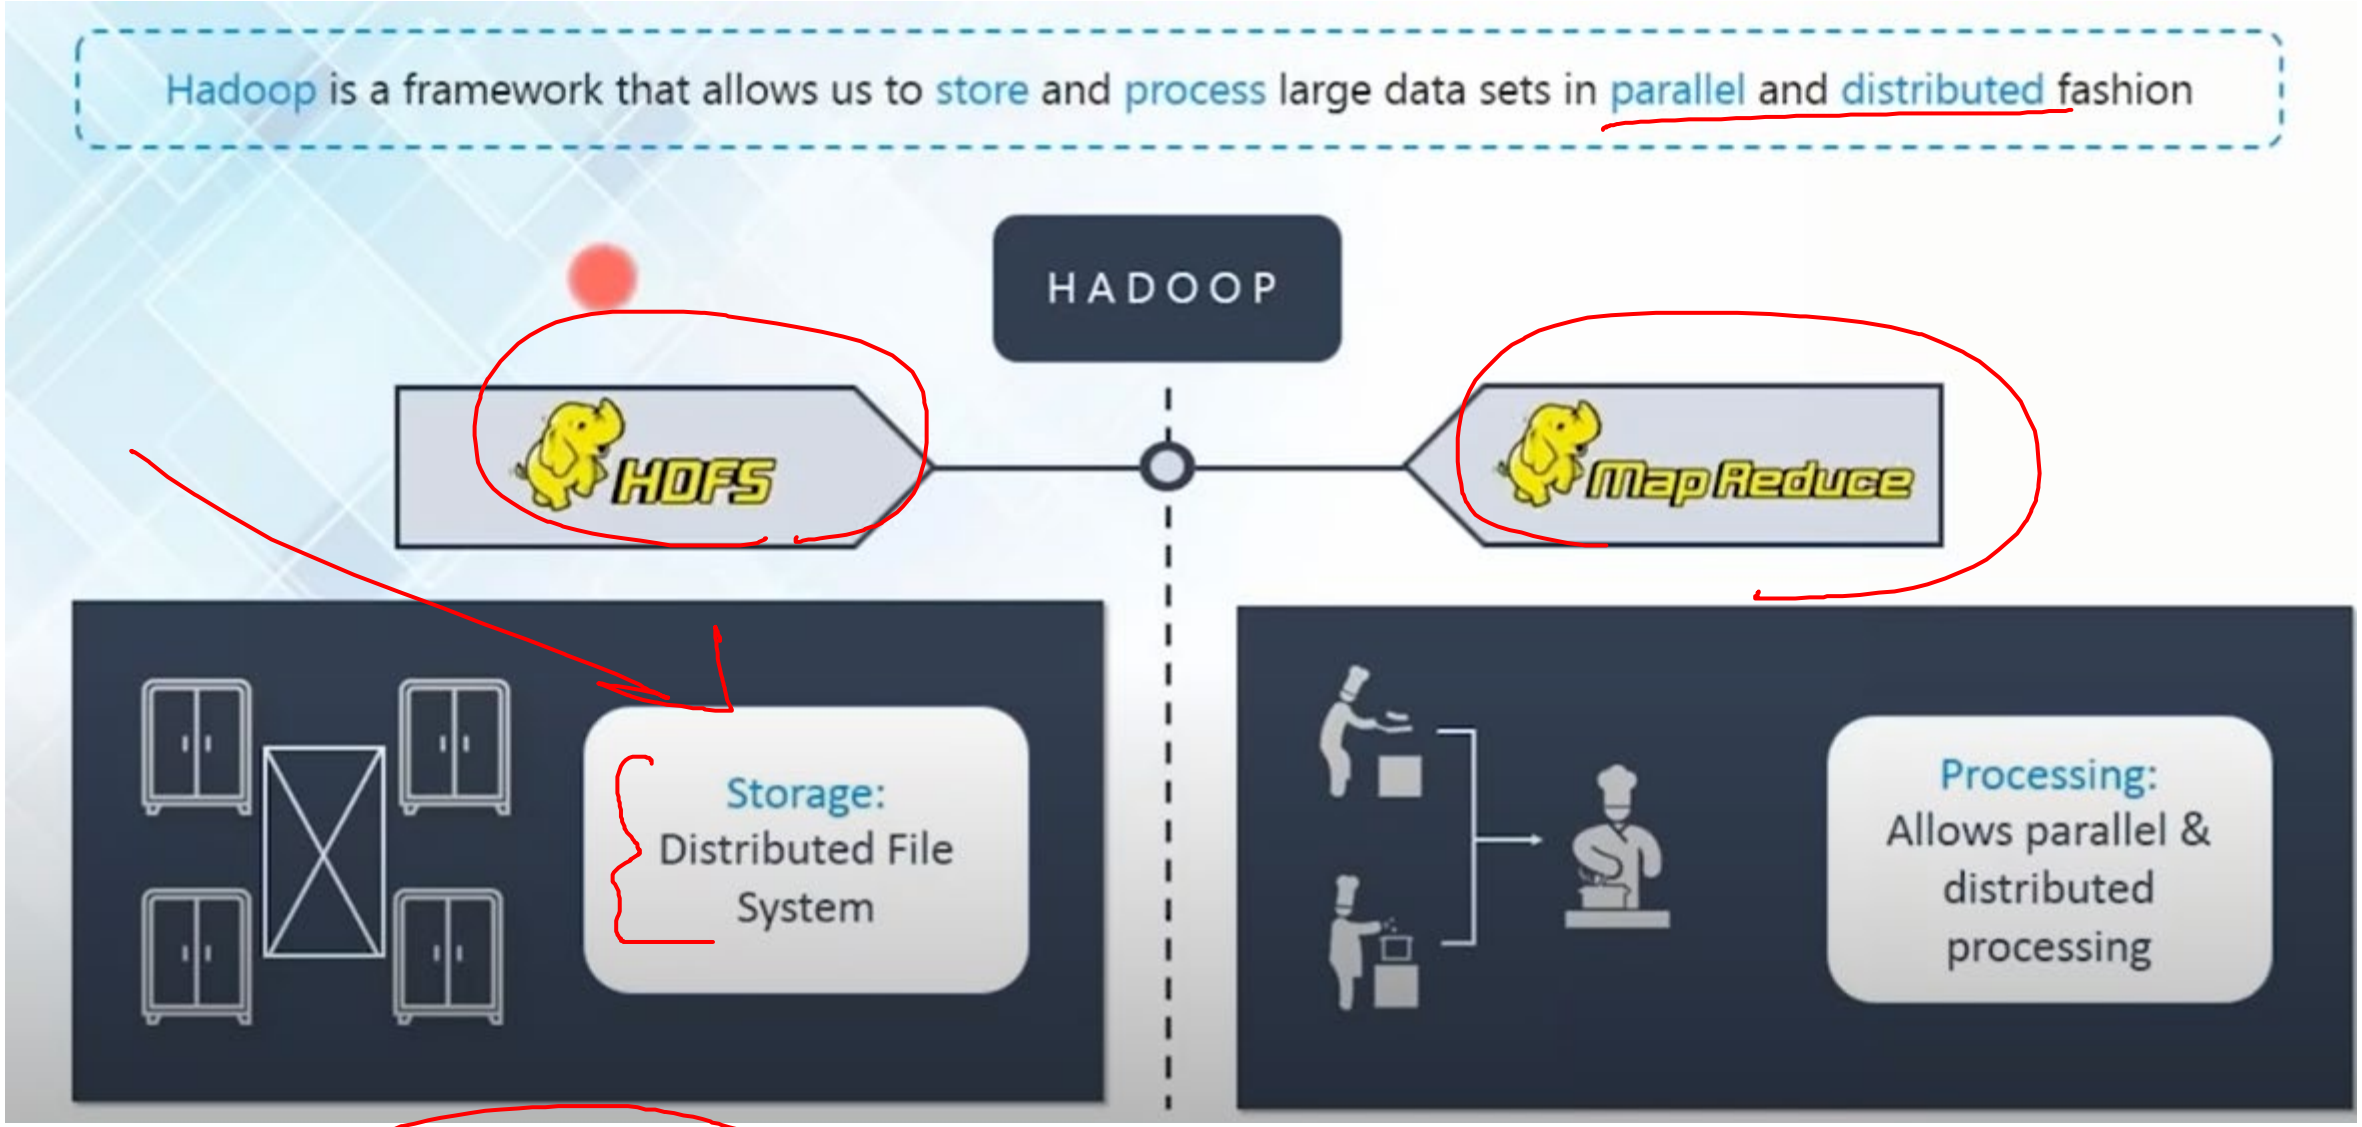

Hadoop Cluster

Hadoop: Master/Slave Architecture

Scenario:

A project Manager managing a team of four employees. He assigns project to each of them and tracks the progress

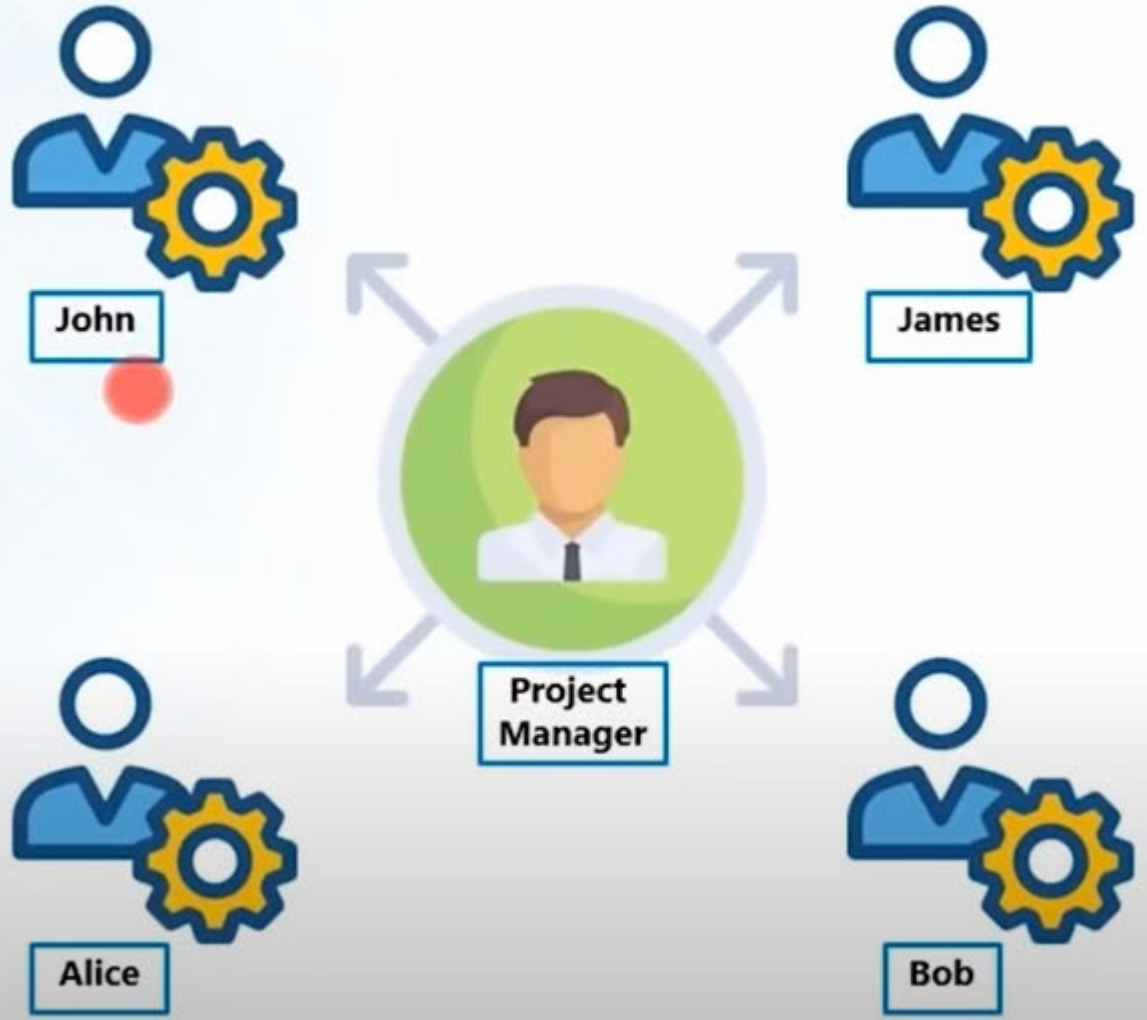

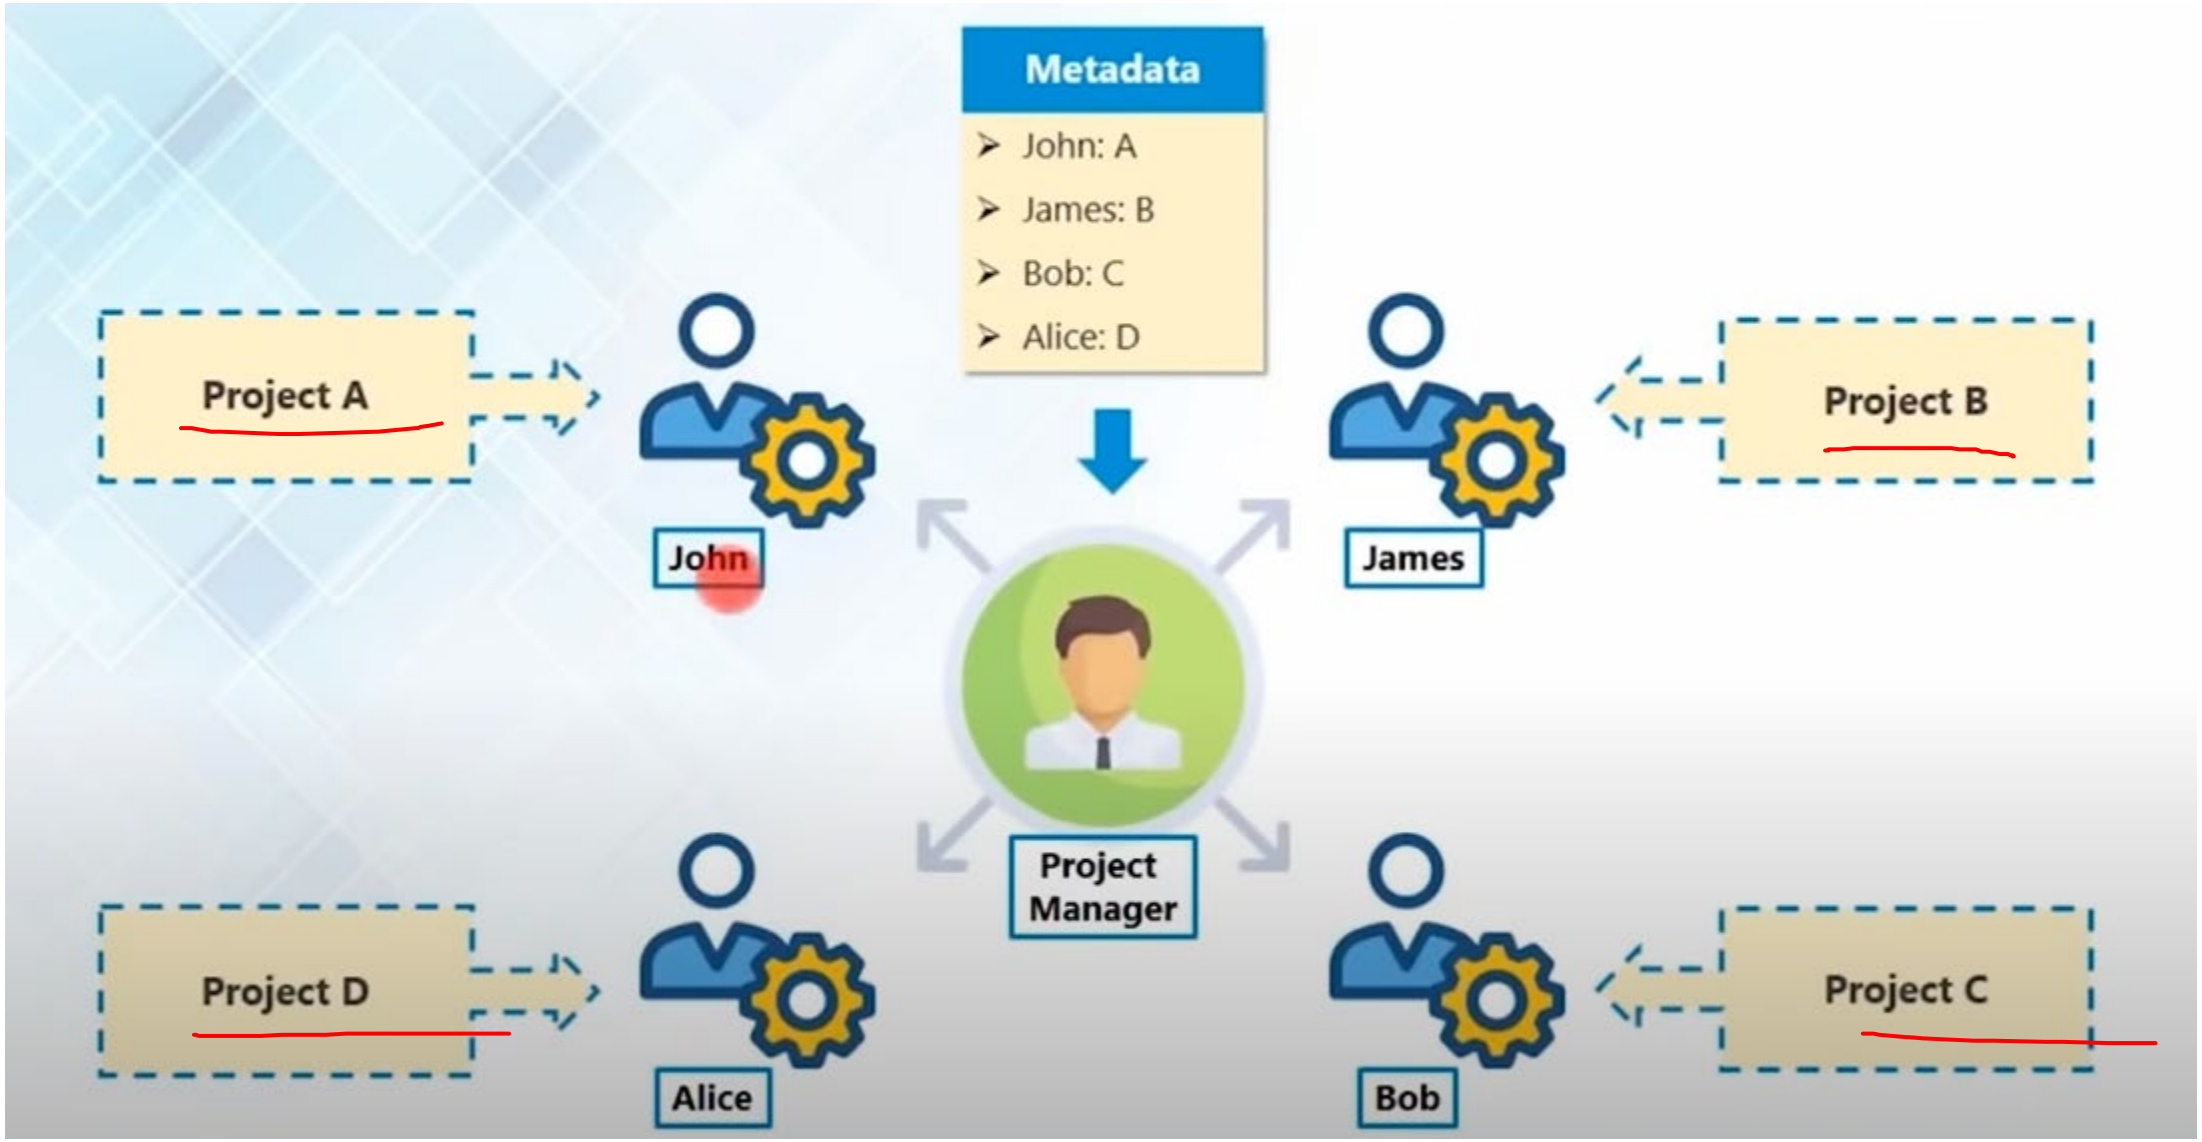
Metadata
John: A
James: B
Bob: C
Alice: D
Project A
Project B
Project D
Project C
John
James
Alice
Bob
Project Manager

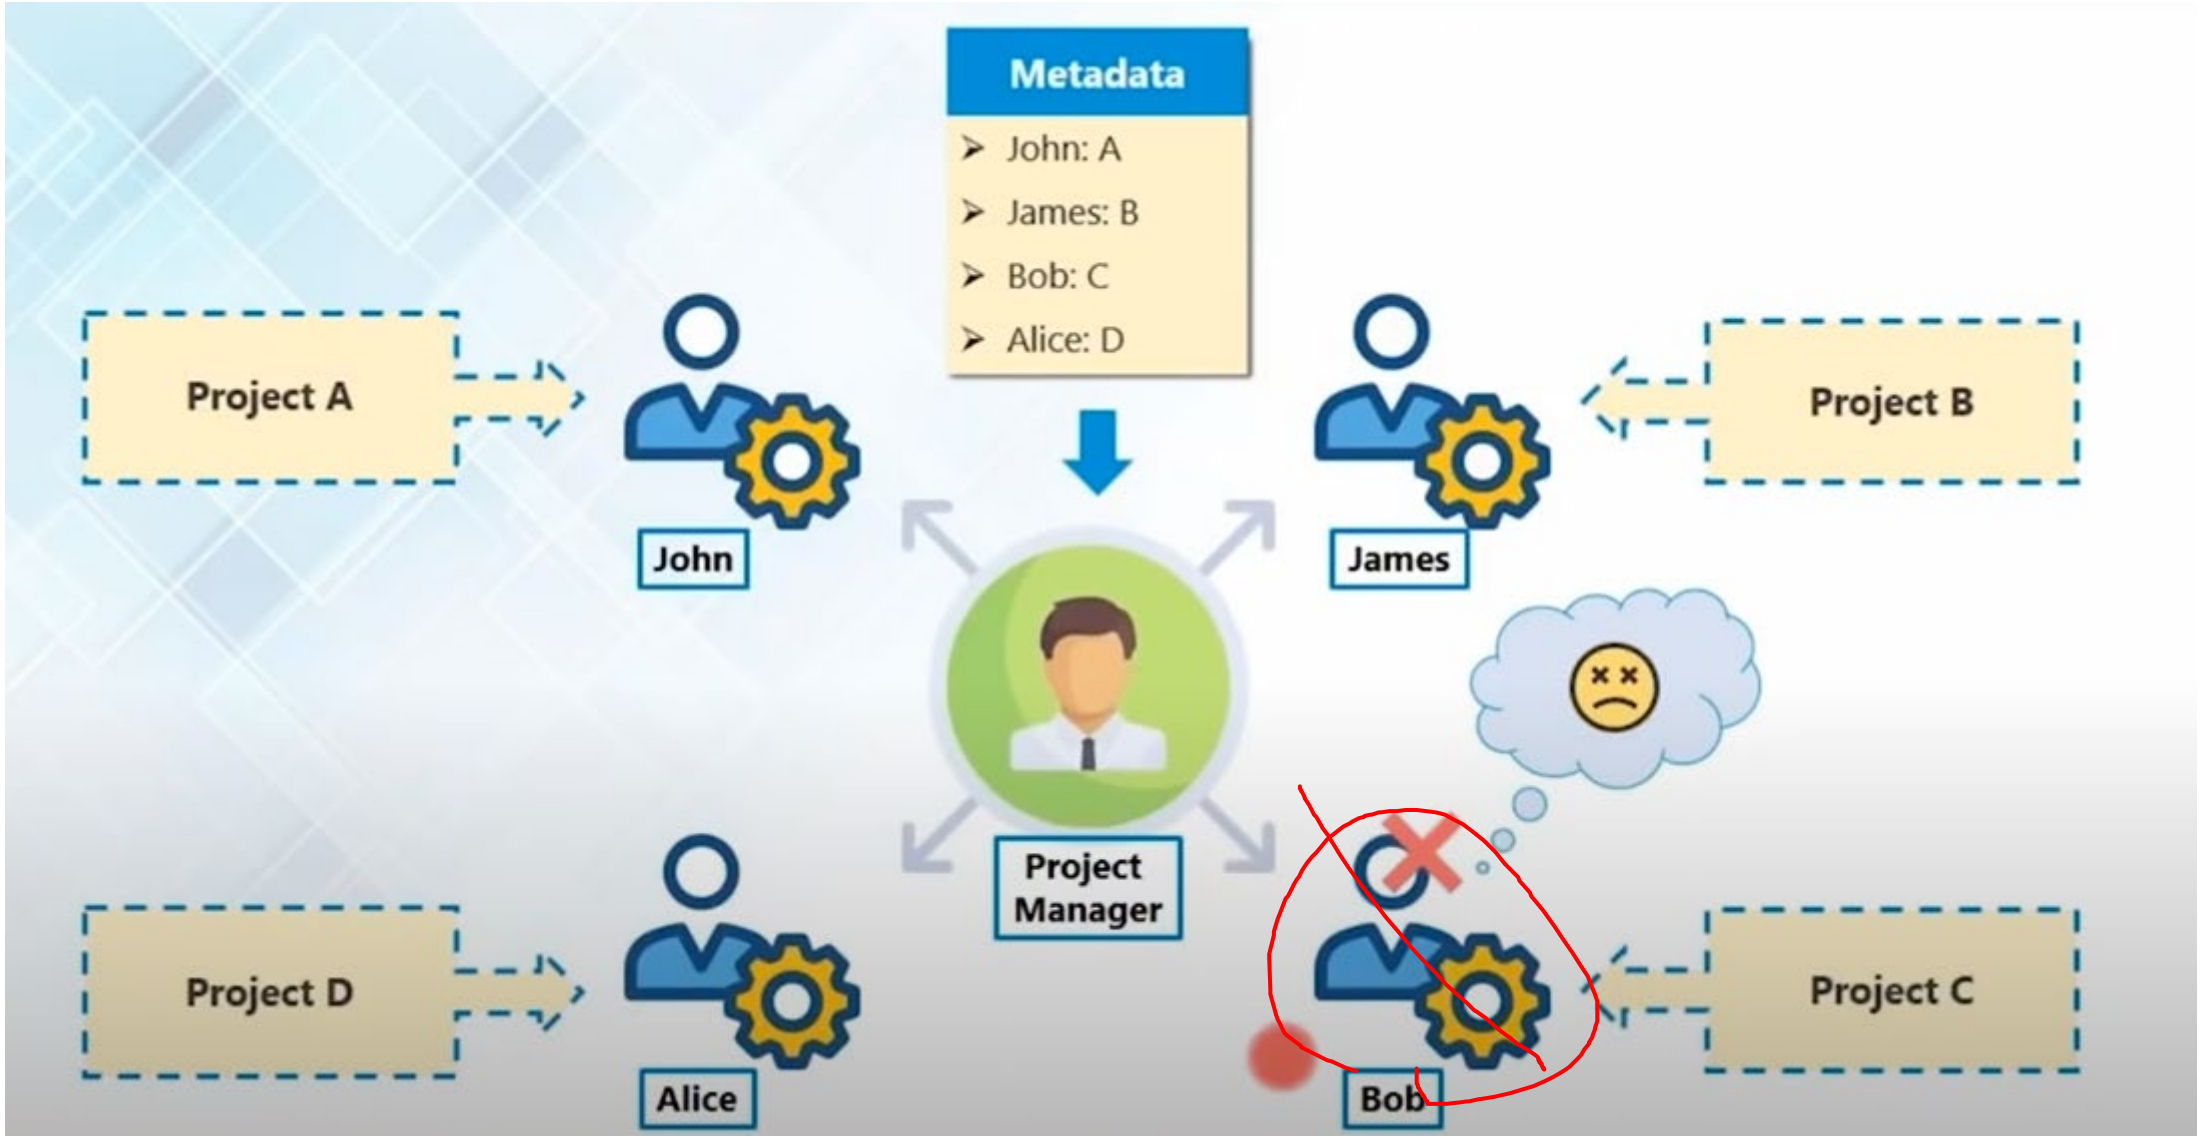
Metadata
John: A
James: B
Bob: C
Alice: D
Project A
Project B
Project C
Project D
John
James
Alice
Bob
Project Manager

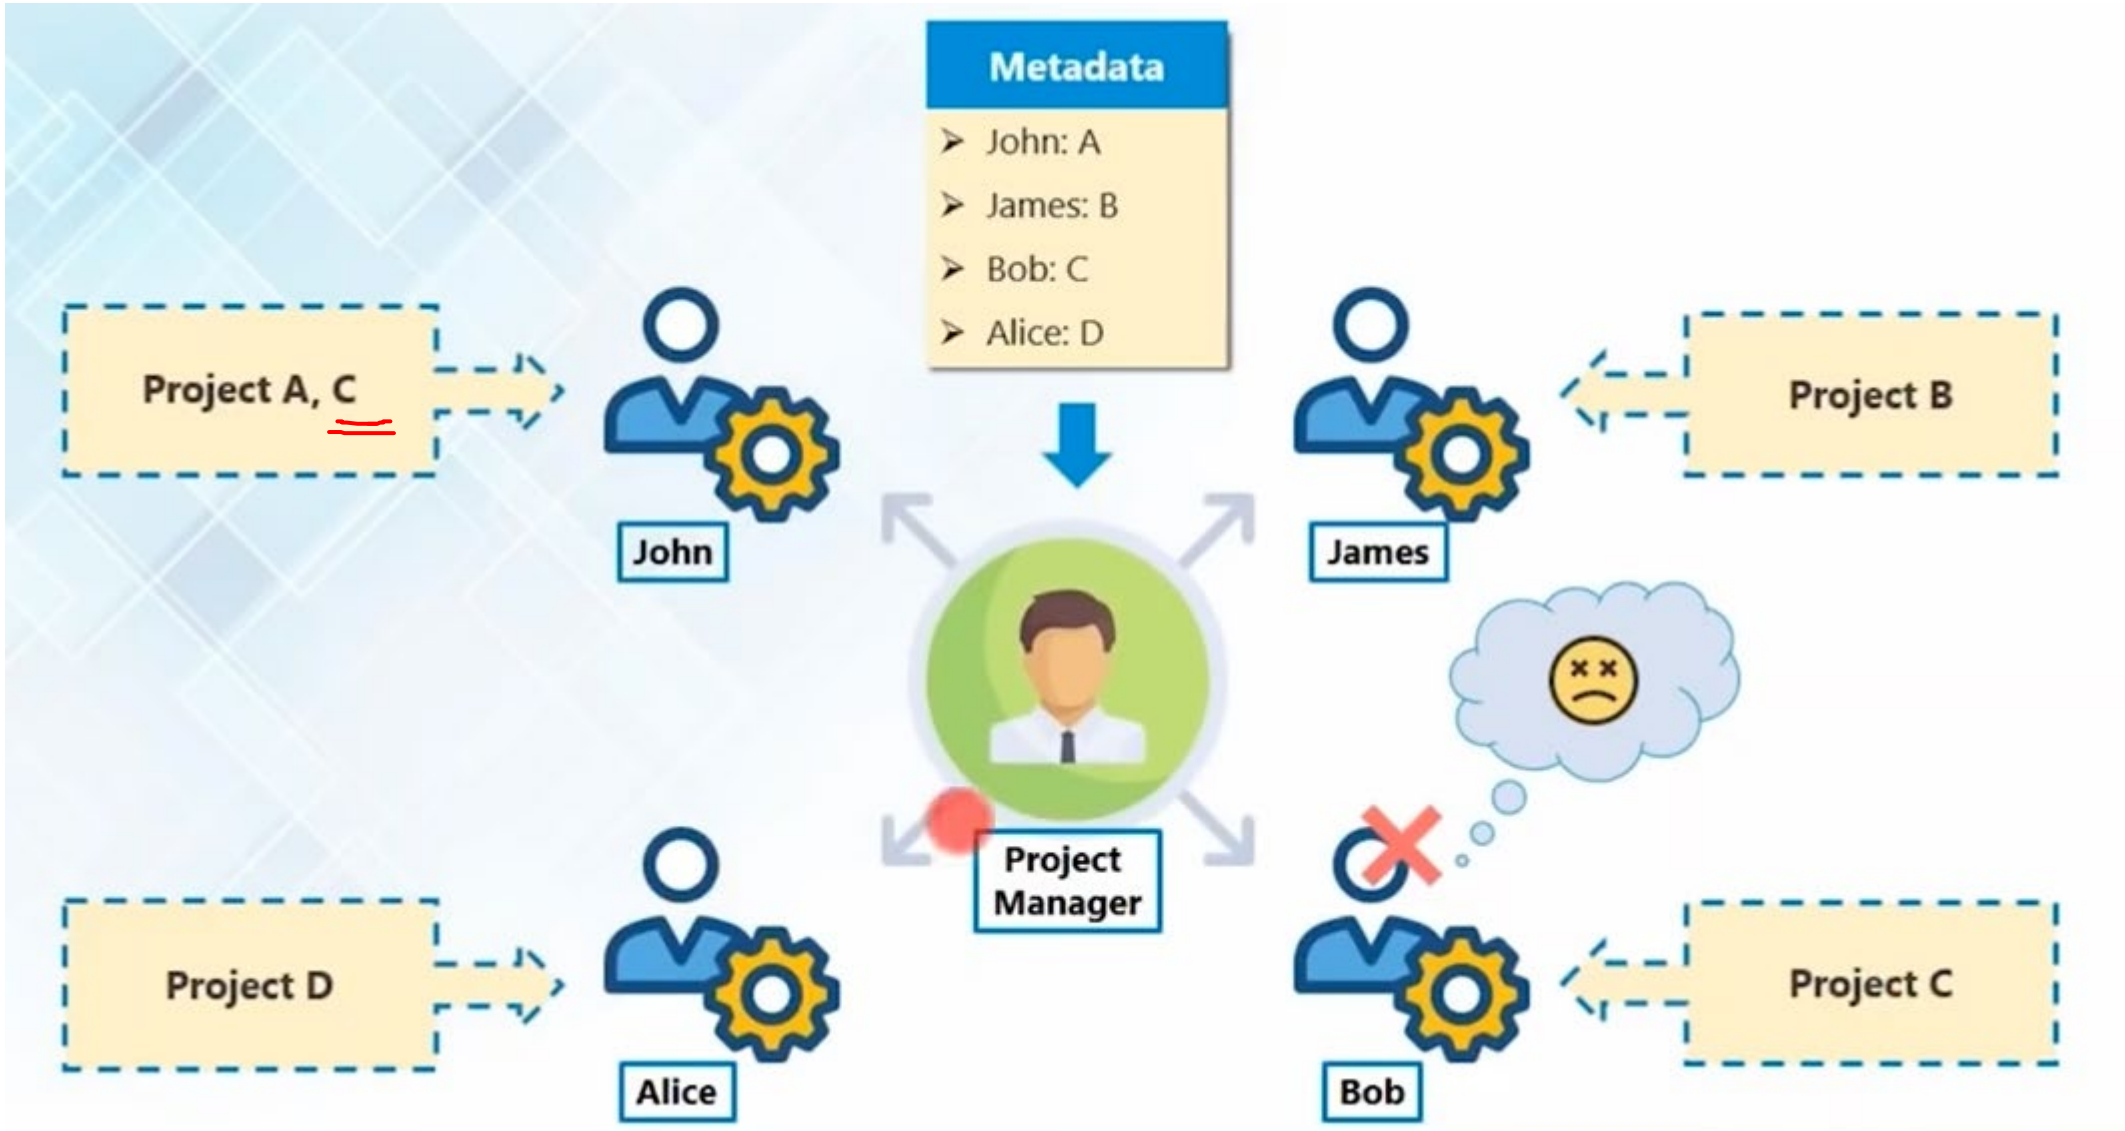
Metadata
John: A
James: B
Bob: C
Alice: D
Project A, C
John
James
Project B
Project Manager
Project D
Alice
Bob
Project C

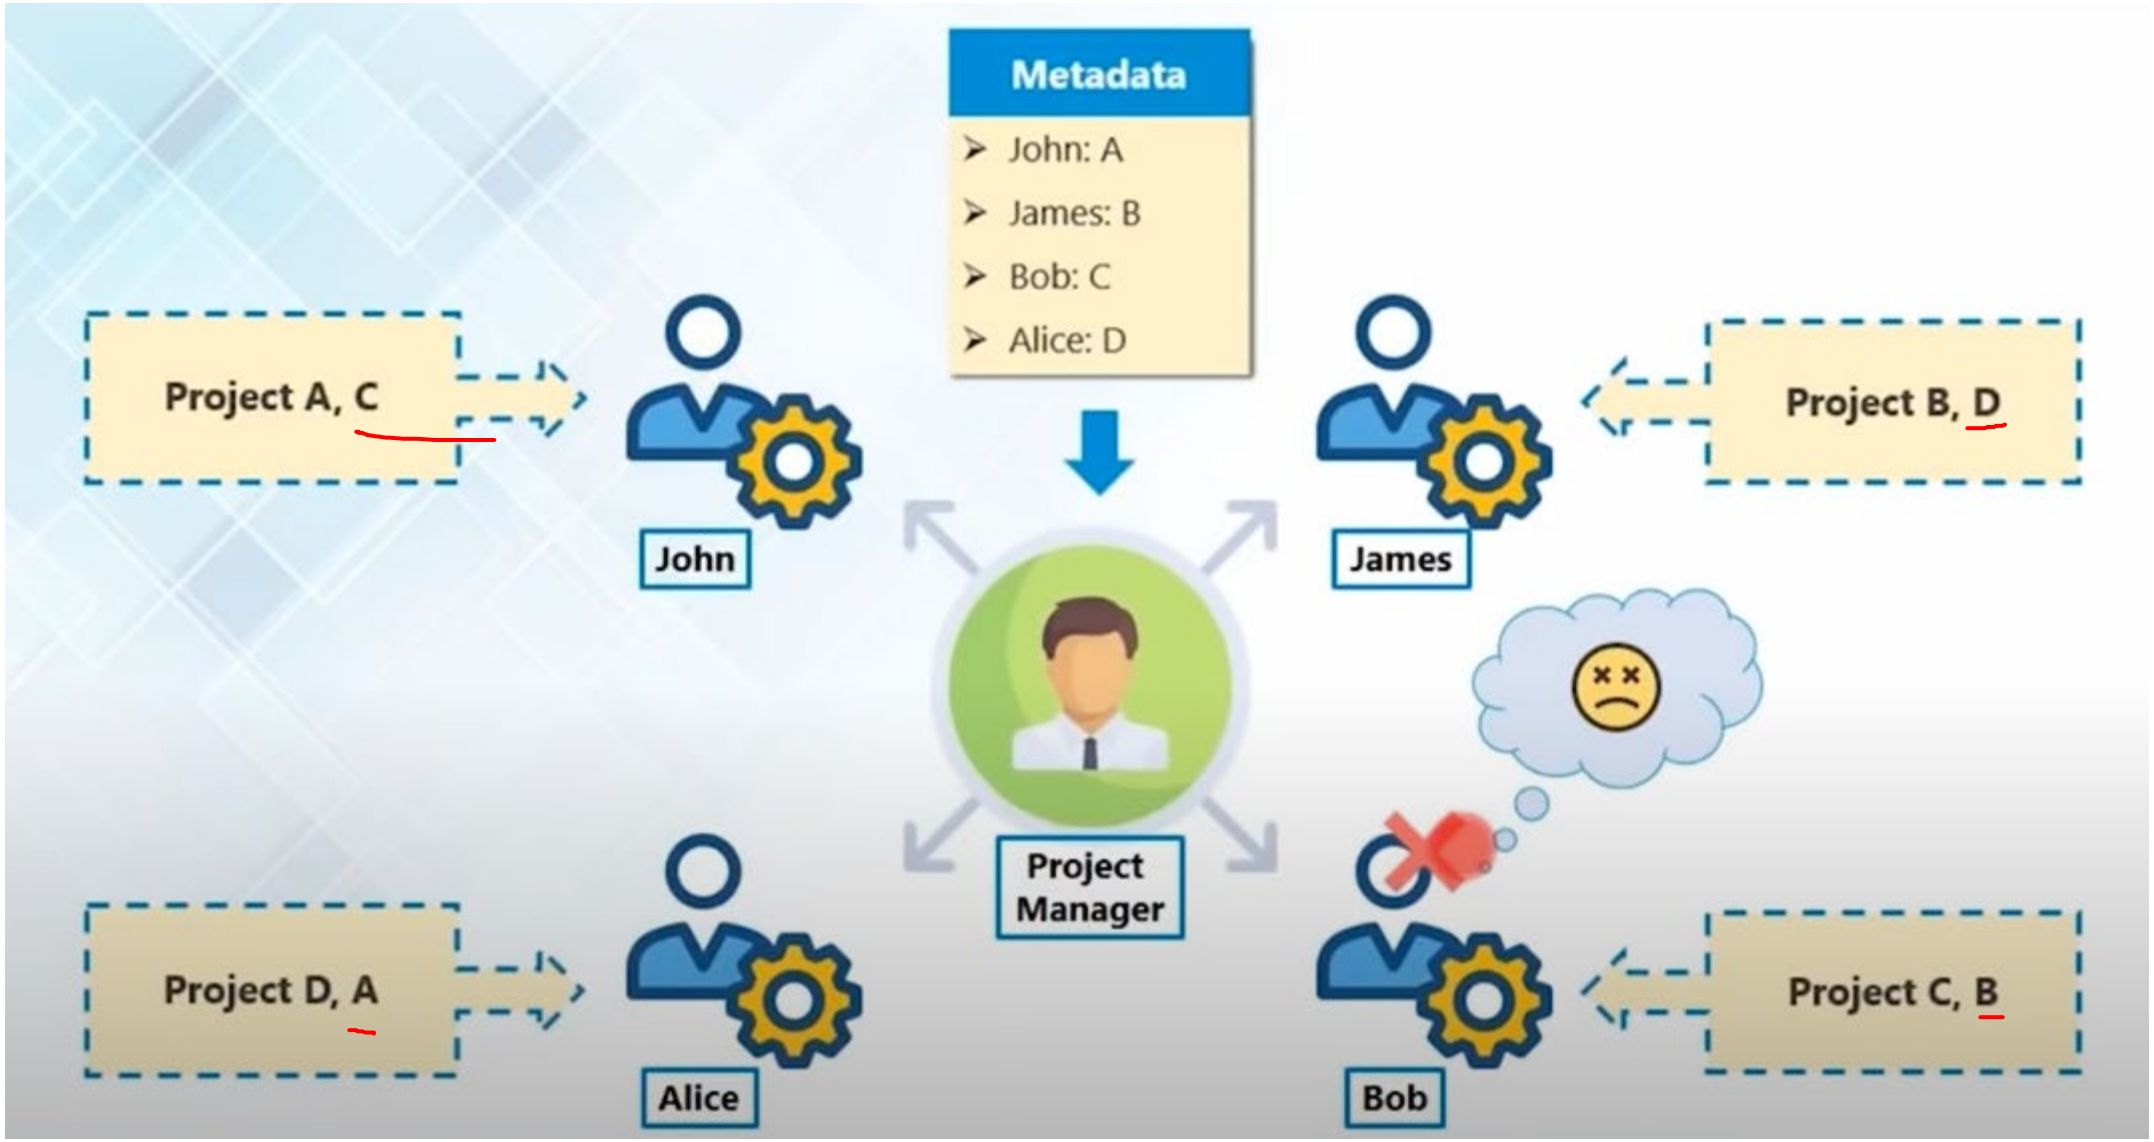
Metadata
John: A
James: B
Bob: C
Alice: D
Project A, C
John
Project B, D
James
Project Manager
Project D, A
Alice
Project C, B
Bob

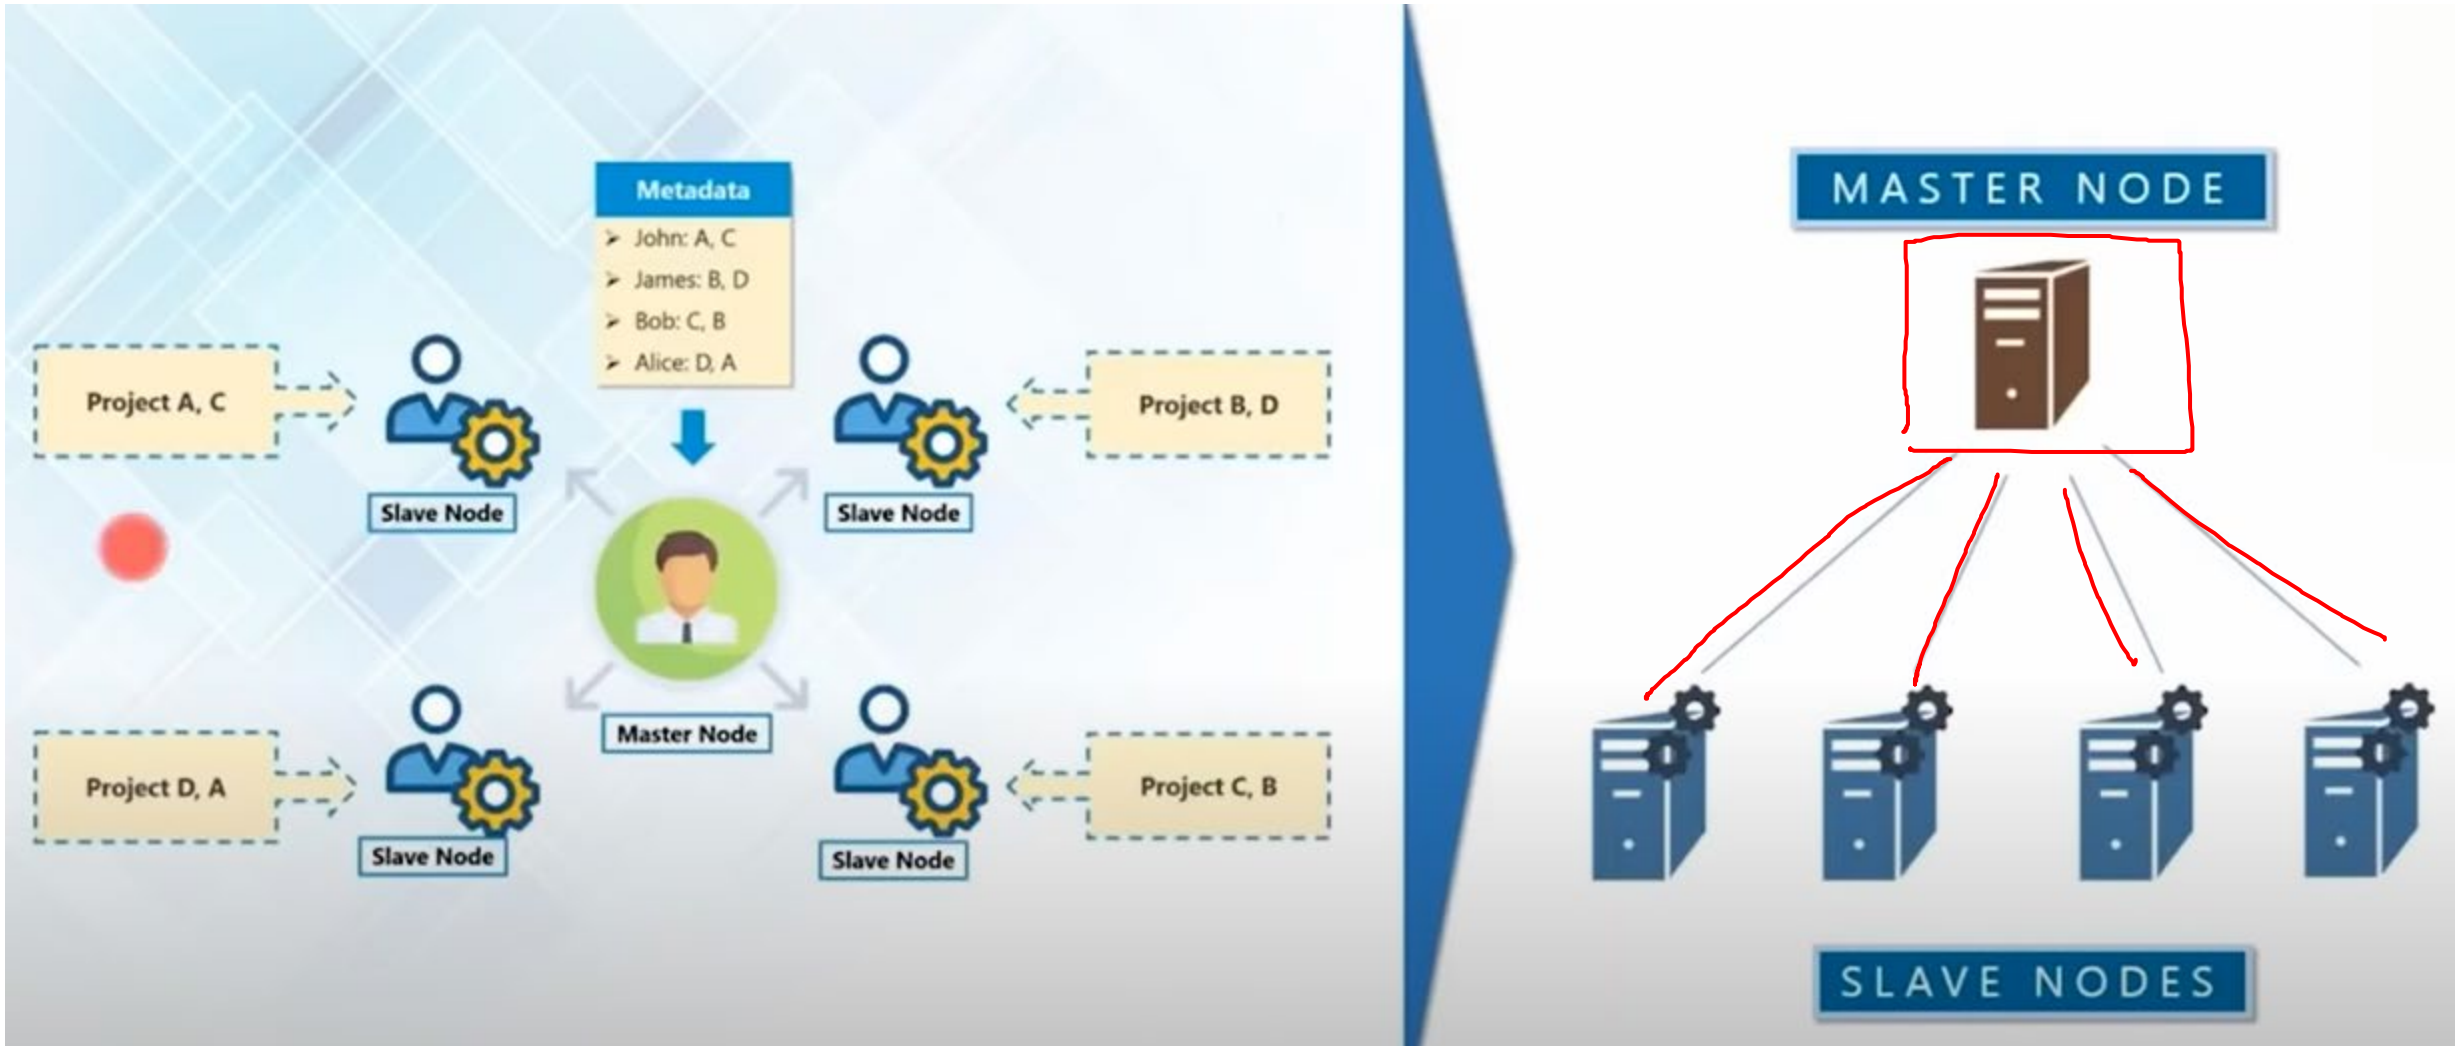
Metadata
John: A, C
James: B, D
Bob: C, B
Alice: D, A
Project A, C
Project B, D
Project D, A
Project C, B
Slave Node
Slave Node
Slave Node
Slave Node
Master Node
MASTER NODE
SLAVE NODES

HADOOP CORE COMPONENTS

HDFS

MapReduce

Storage:
Distributed File System

Processing:
Allows parallel & distributed processing

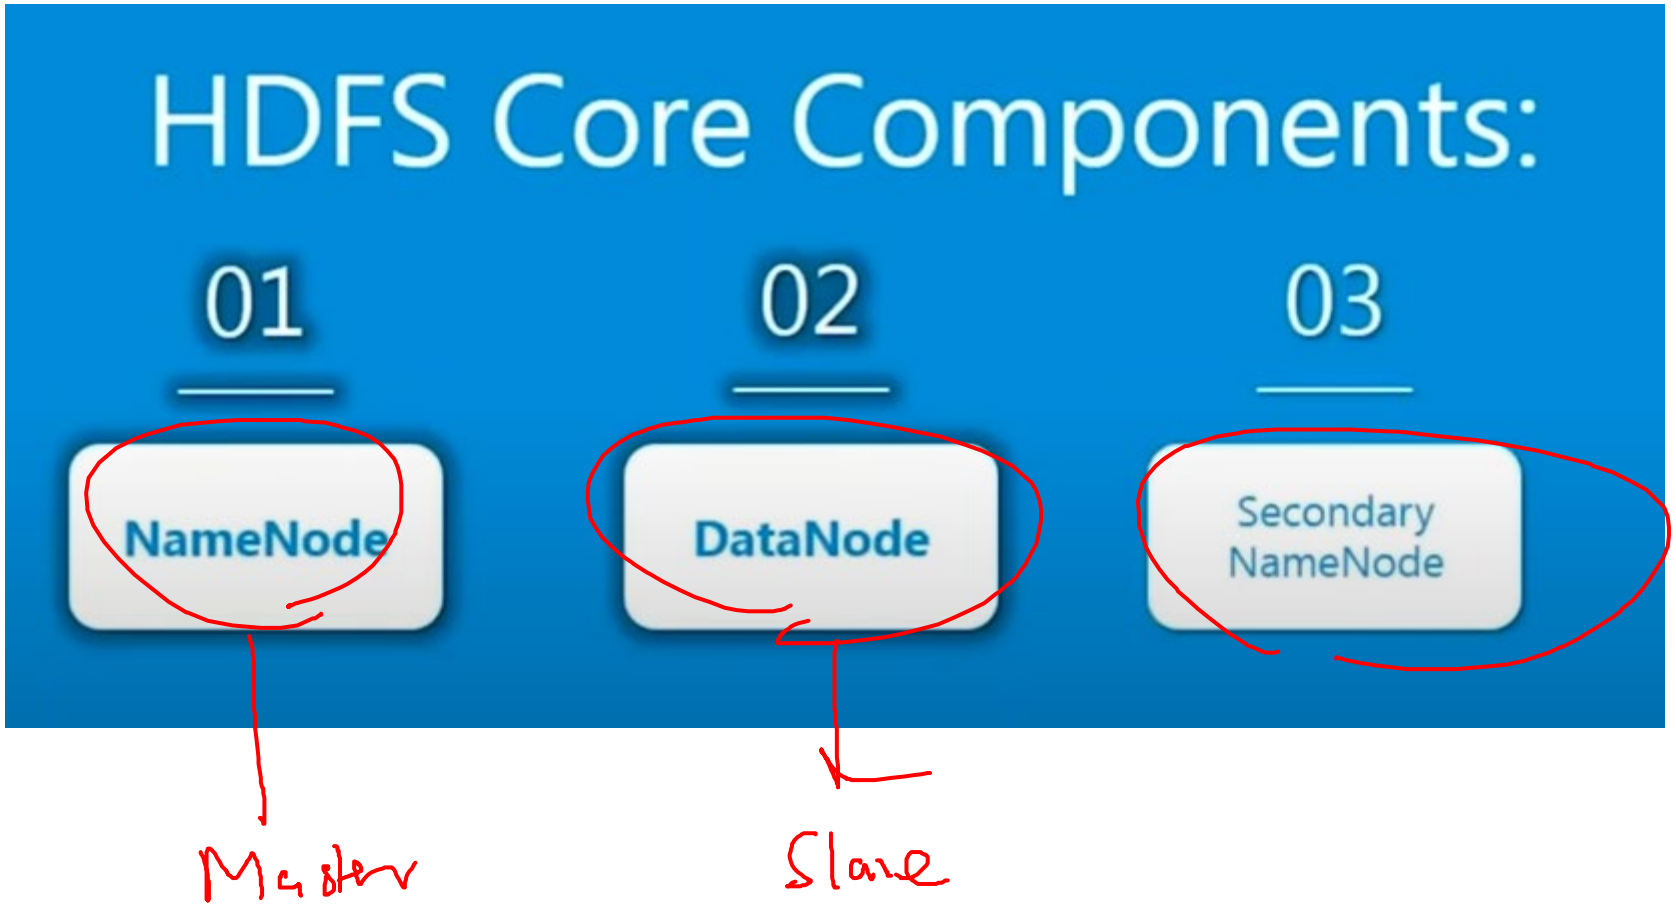
HDFS Core Components:
01
02
03
NameNode
DataNode
Secondary NameNode
Master
Slave

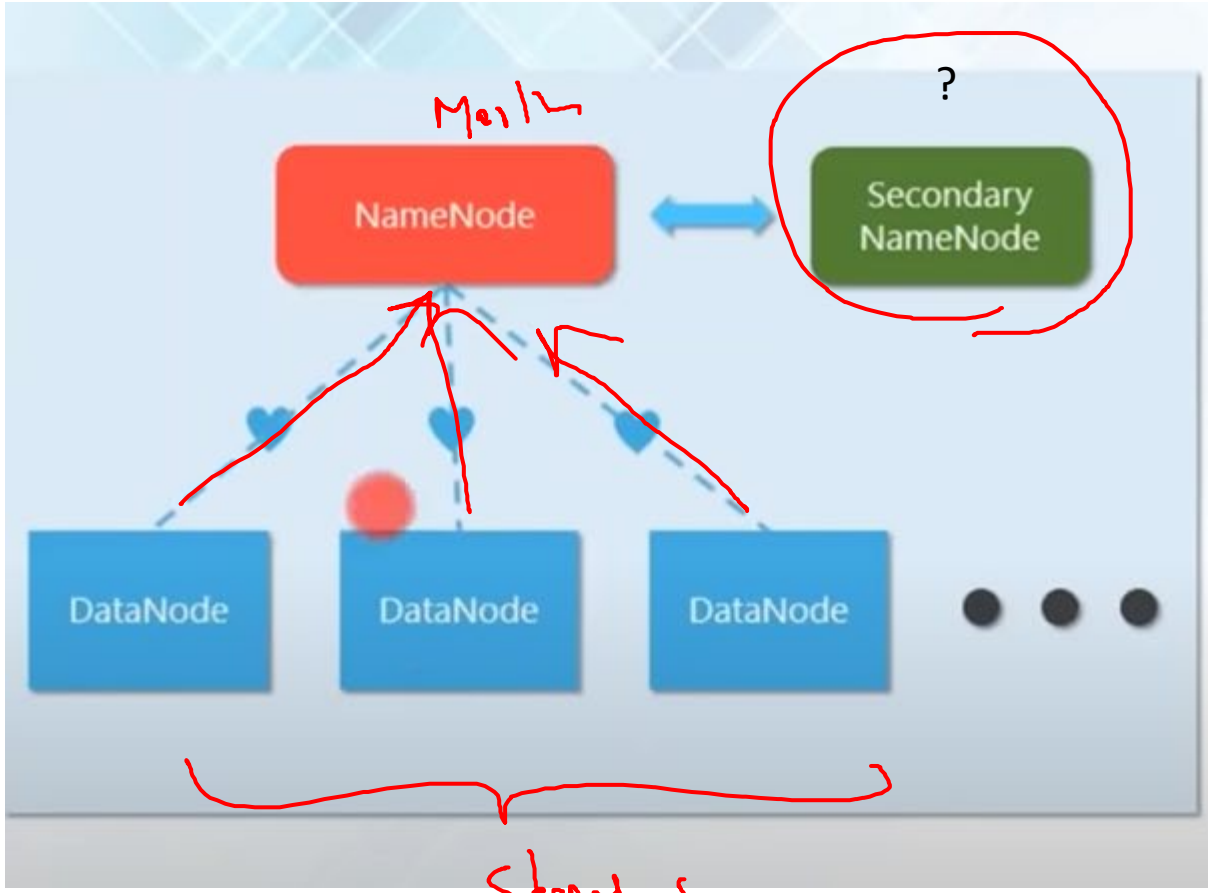

NameNode:

- Maintains and Manages DataNodes
- Records metadata i.e. information about data blocks e.g. location of blocks stored, the size of the files, permissions, hierarchy, etc.
- Receives heartbeat and block report from all the DataNodes

DataNode:

- Slave daemons
- Stores actual data
- Serves read and write requests from the clients

Meta data = changes made to a file

- Checkpointing is a process of combining edit logs with FsImage
- Secondary NameNode takes over the responsibility of checkpointing, therefore, making NameNode more available
- Allows faster Failover as it prevents edit logs from getting too huge
- Checkpointing happens periodically (default: 1 hour)

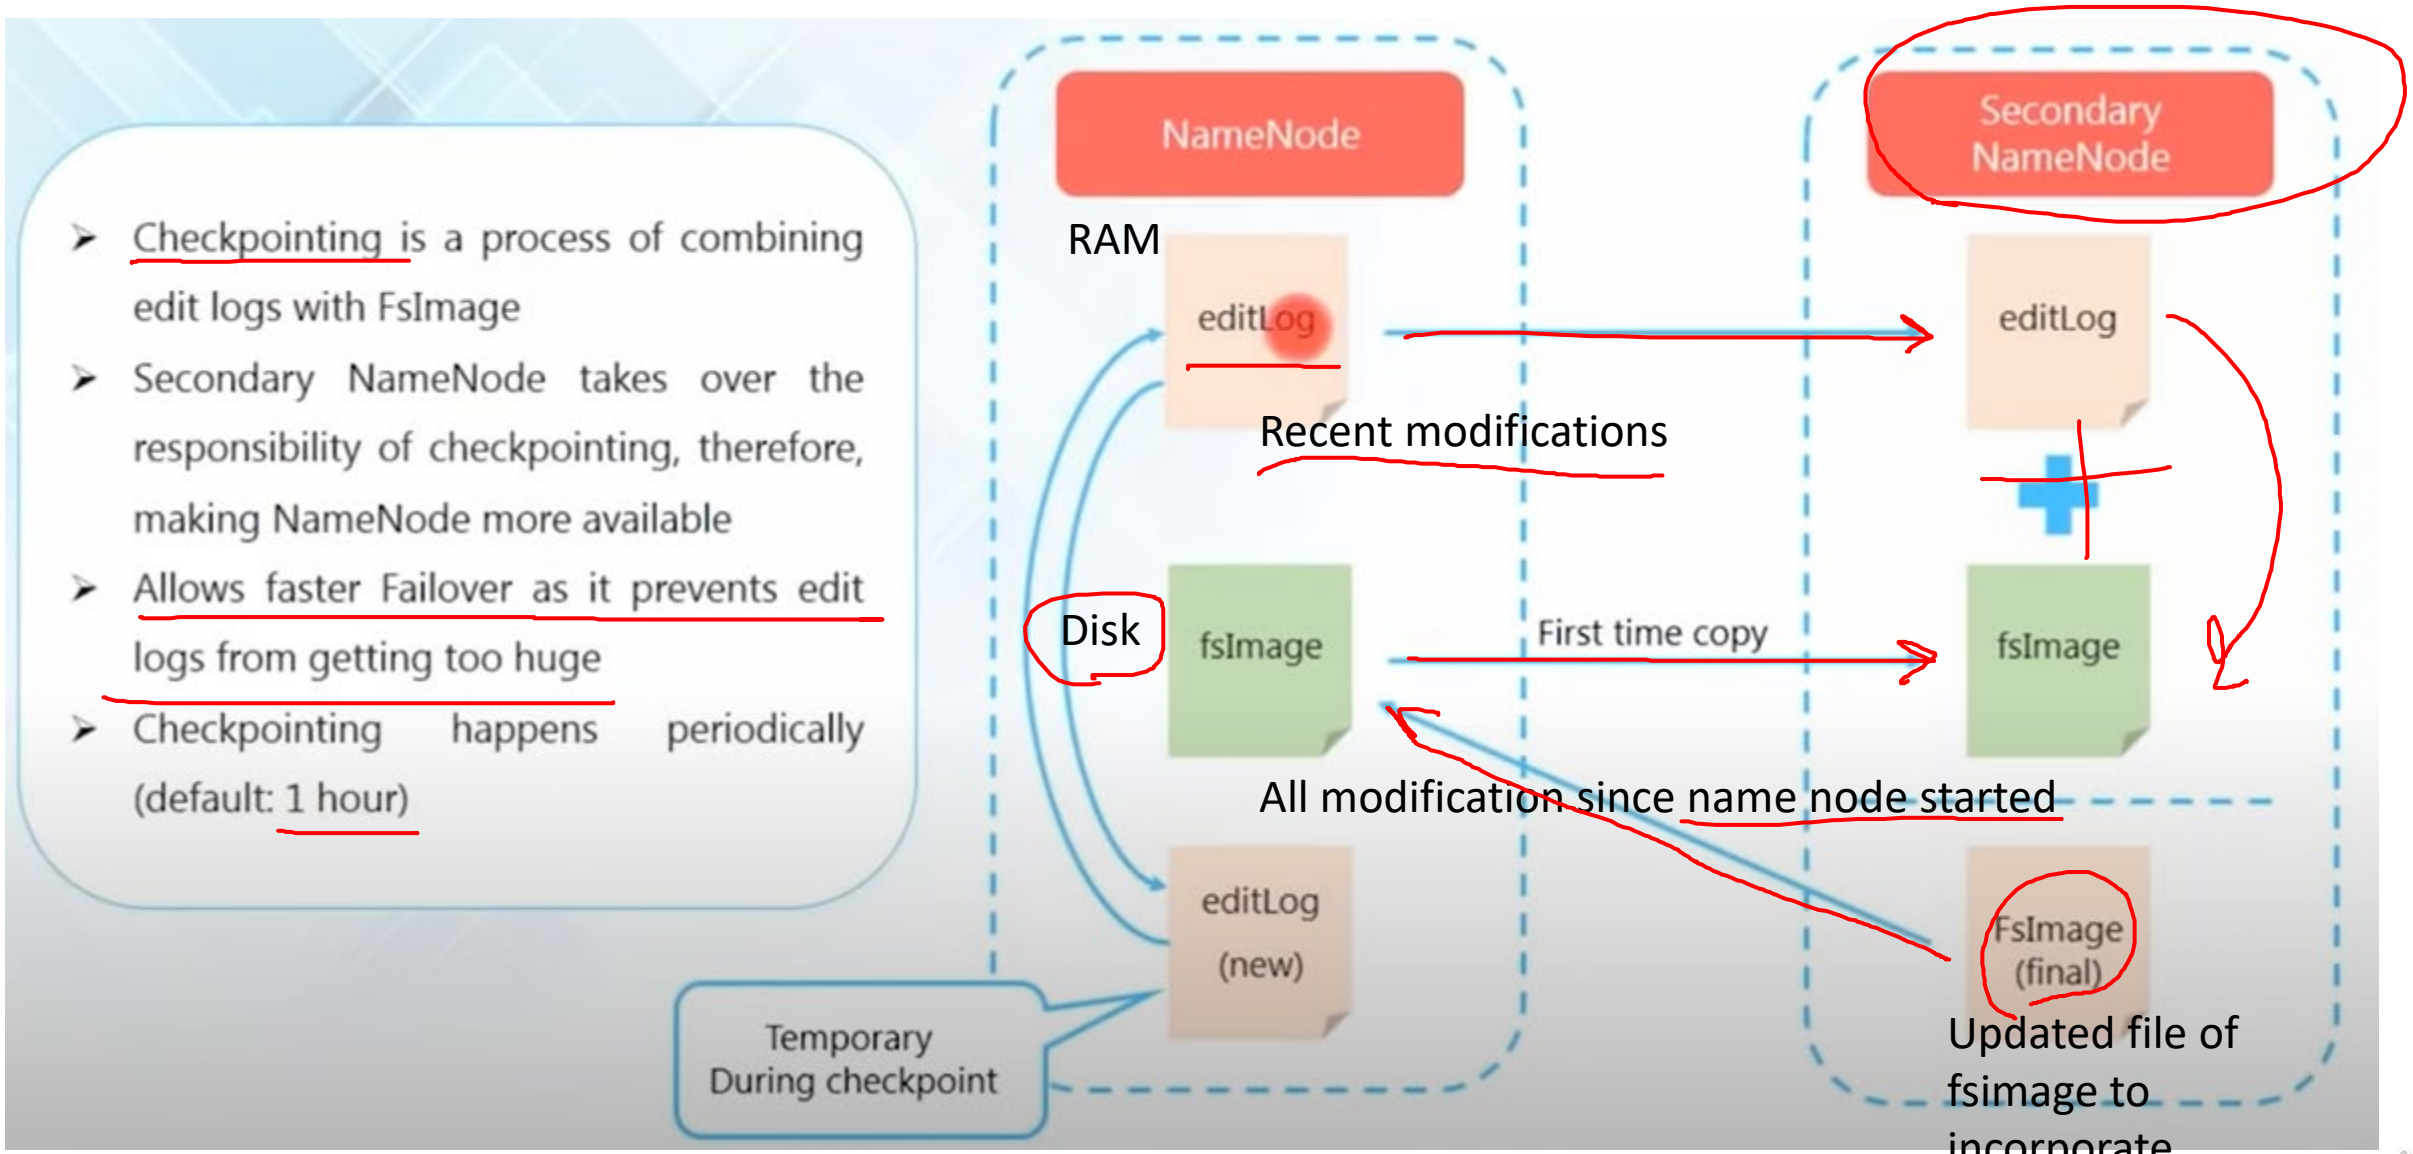

How the data is actually stored in DataNodes?
HDFS Data Blocks

- Each file is stored on HDFS as blocks
- The default size of each block is 128 MB in Apache Hadoop 2.x (64 MB in Apache Hadoop 1.x)

Scalable storage and efficient parallel processing

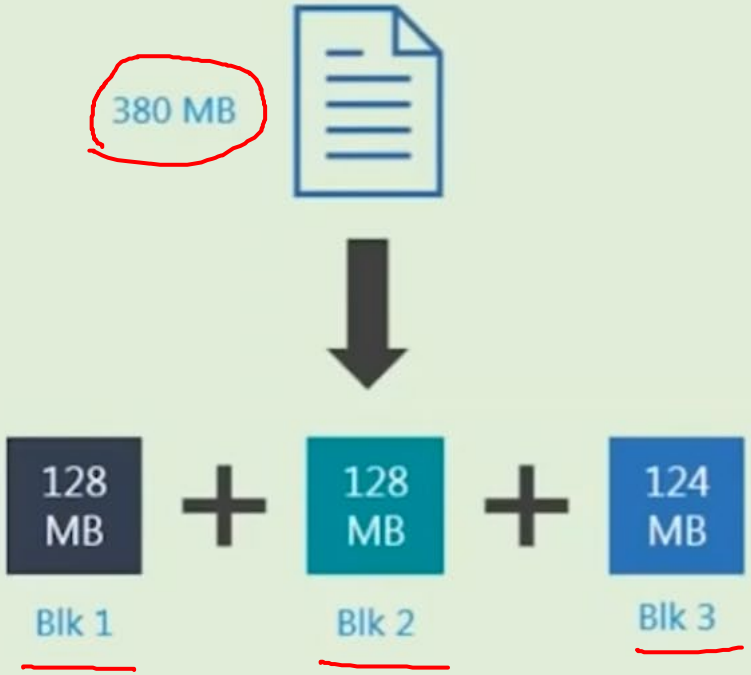

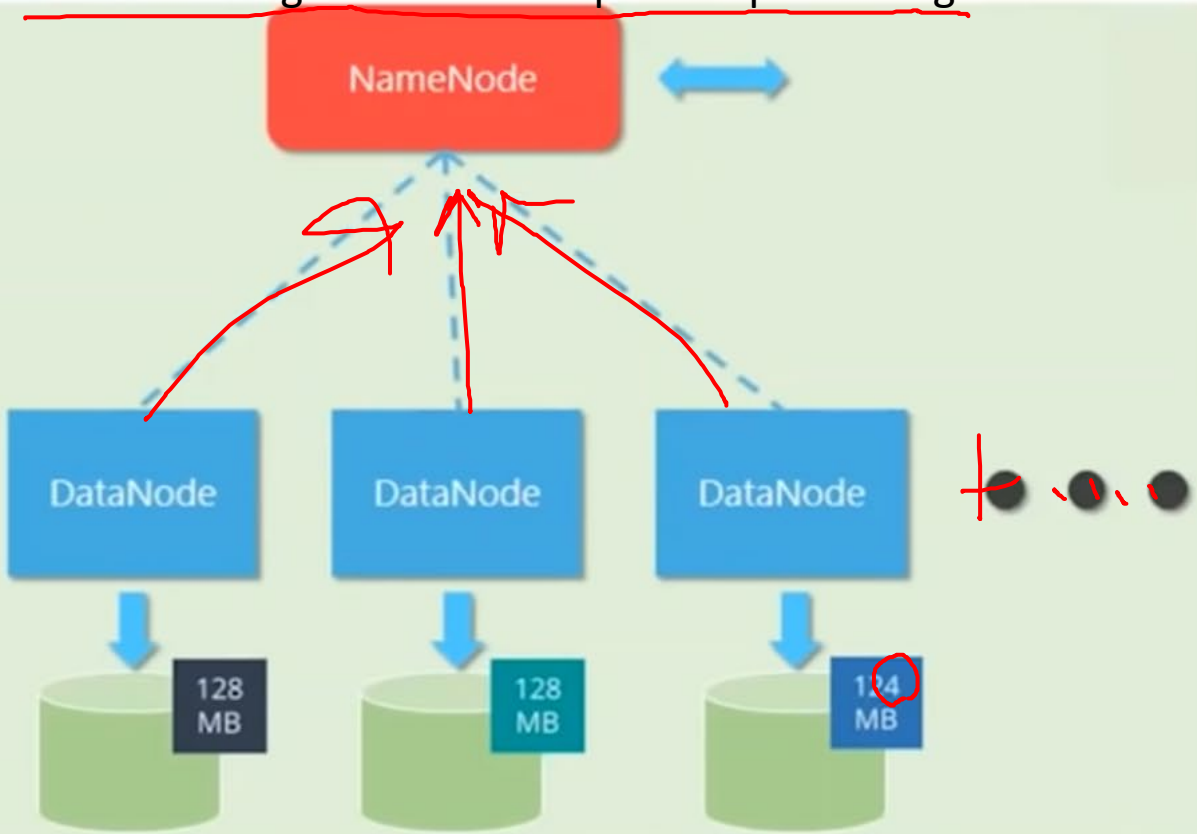

4 MB saved using 128 mb blocks

Fault Tolerance: How Hadoop cope up with DataNode Failure?

Scenario:

One of the DataNodes crashed containing the data blocks

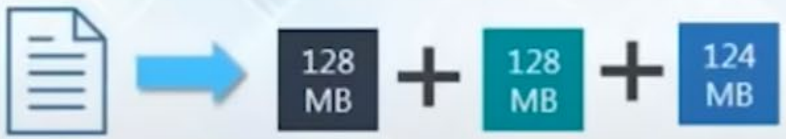

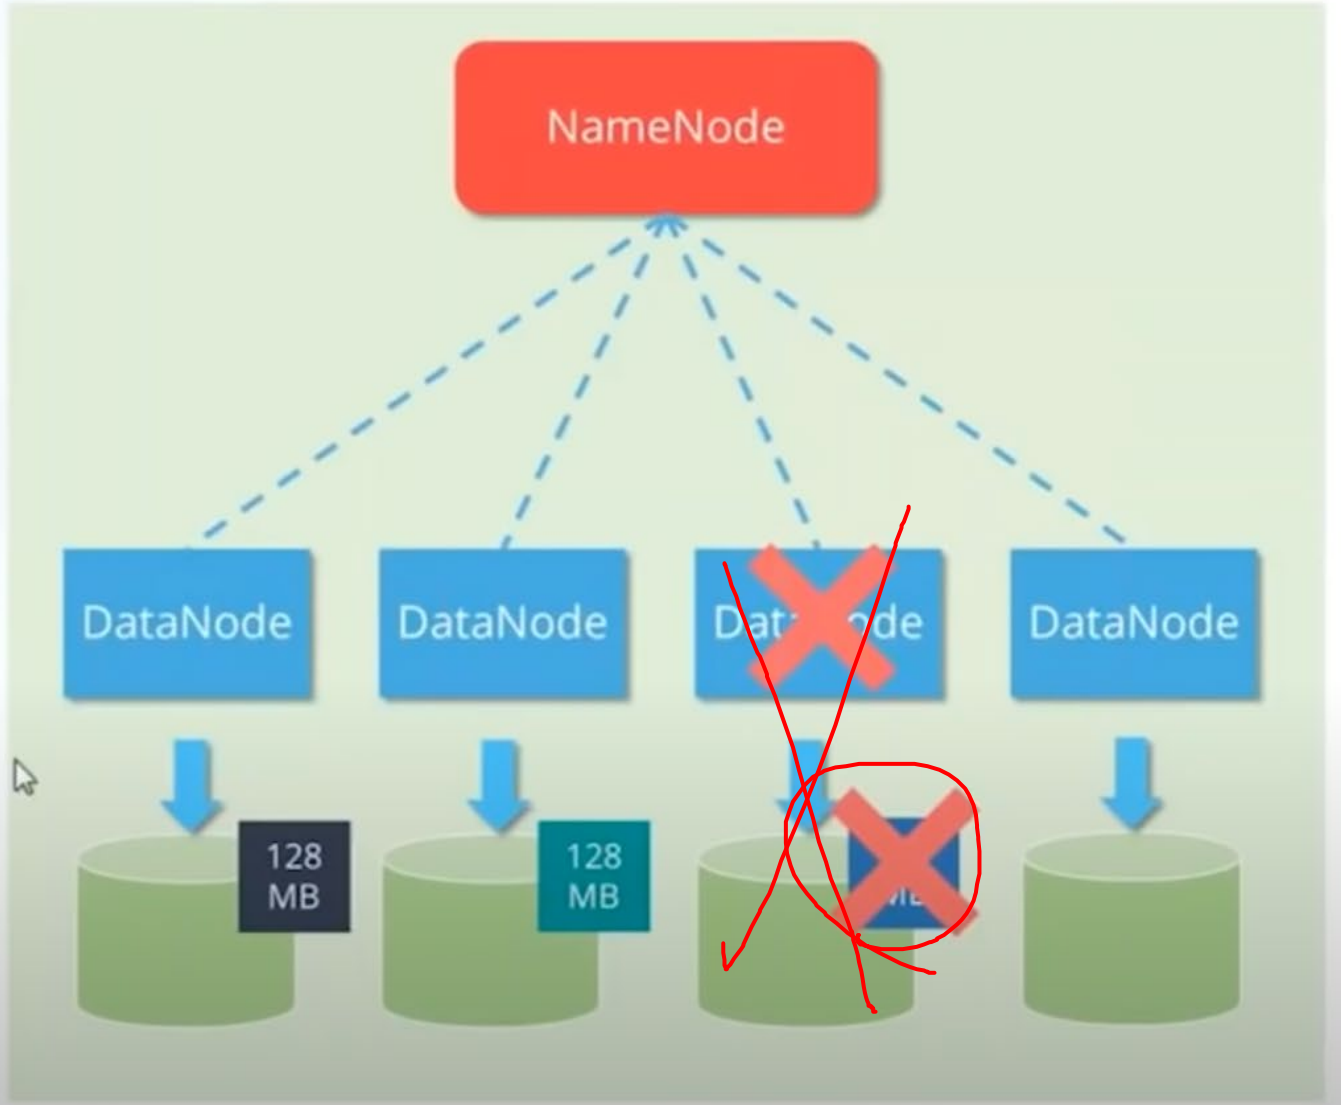

Solution: Replication Factor

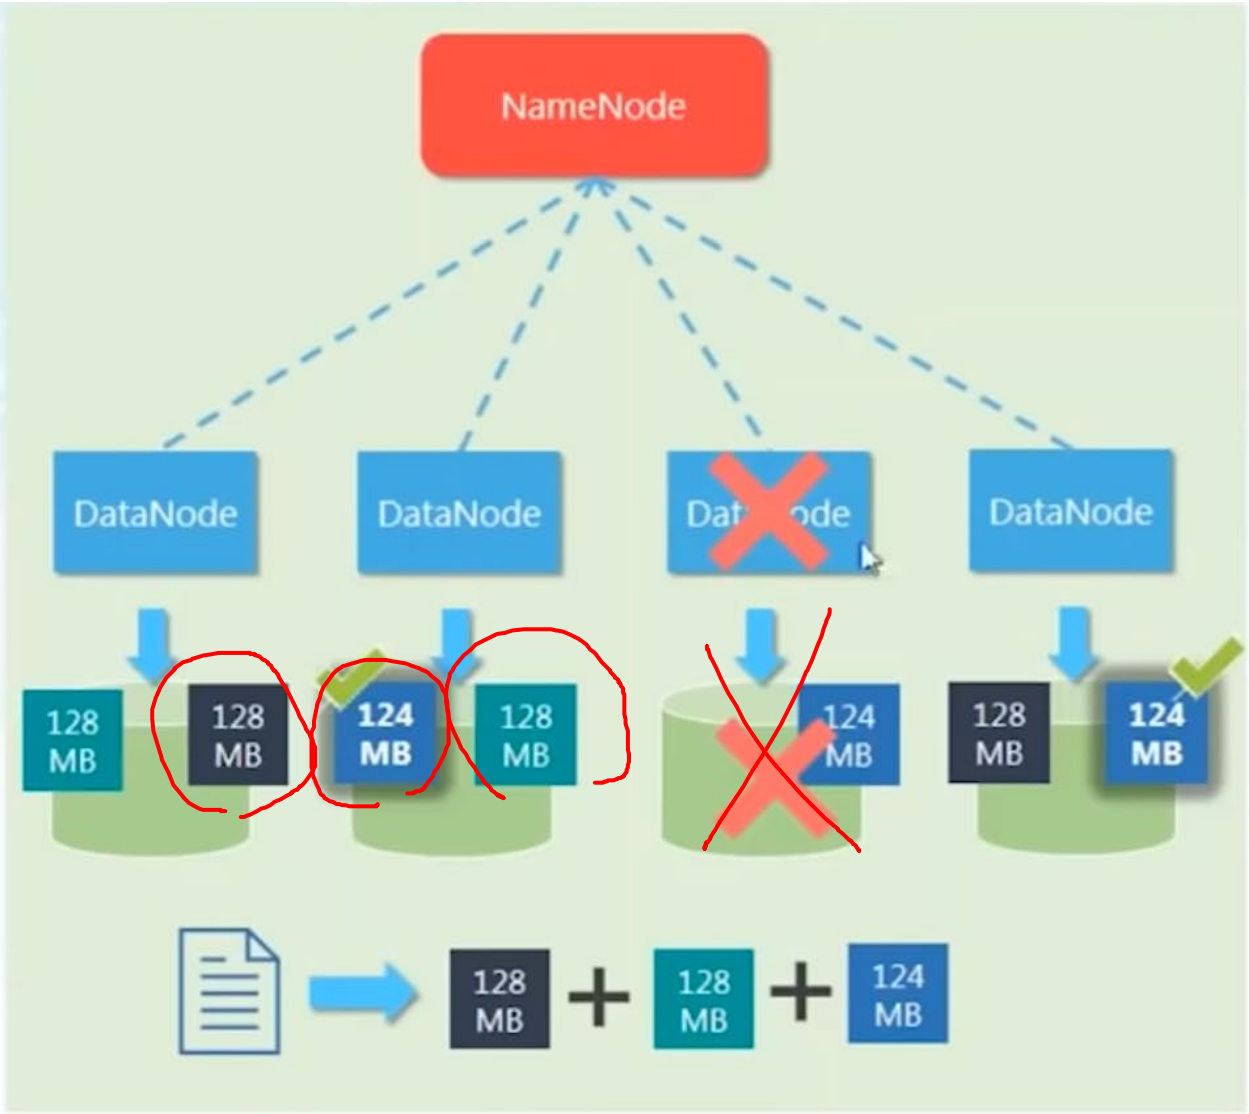

Solution:

Each data blocks are replicated (thrice by default) and are distributed across different DataNodes

Setting up HDFS - Write Pipeline

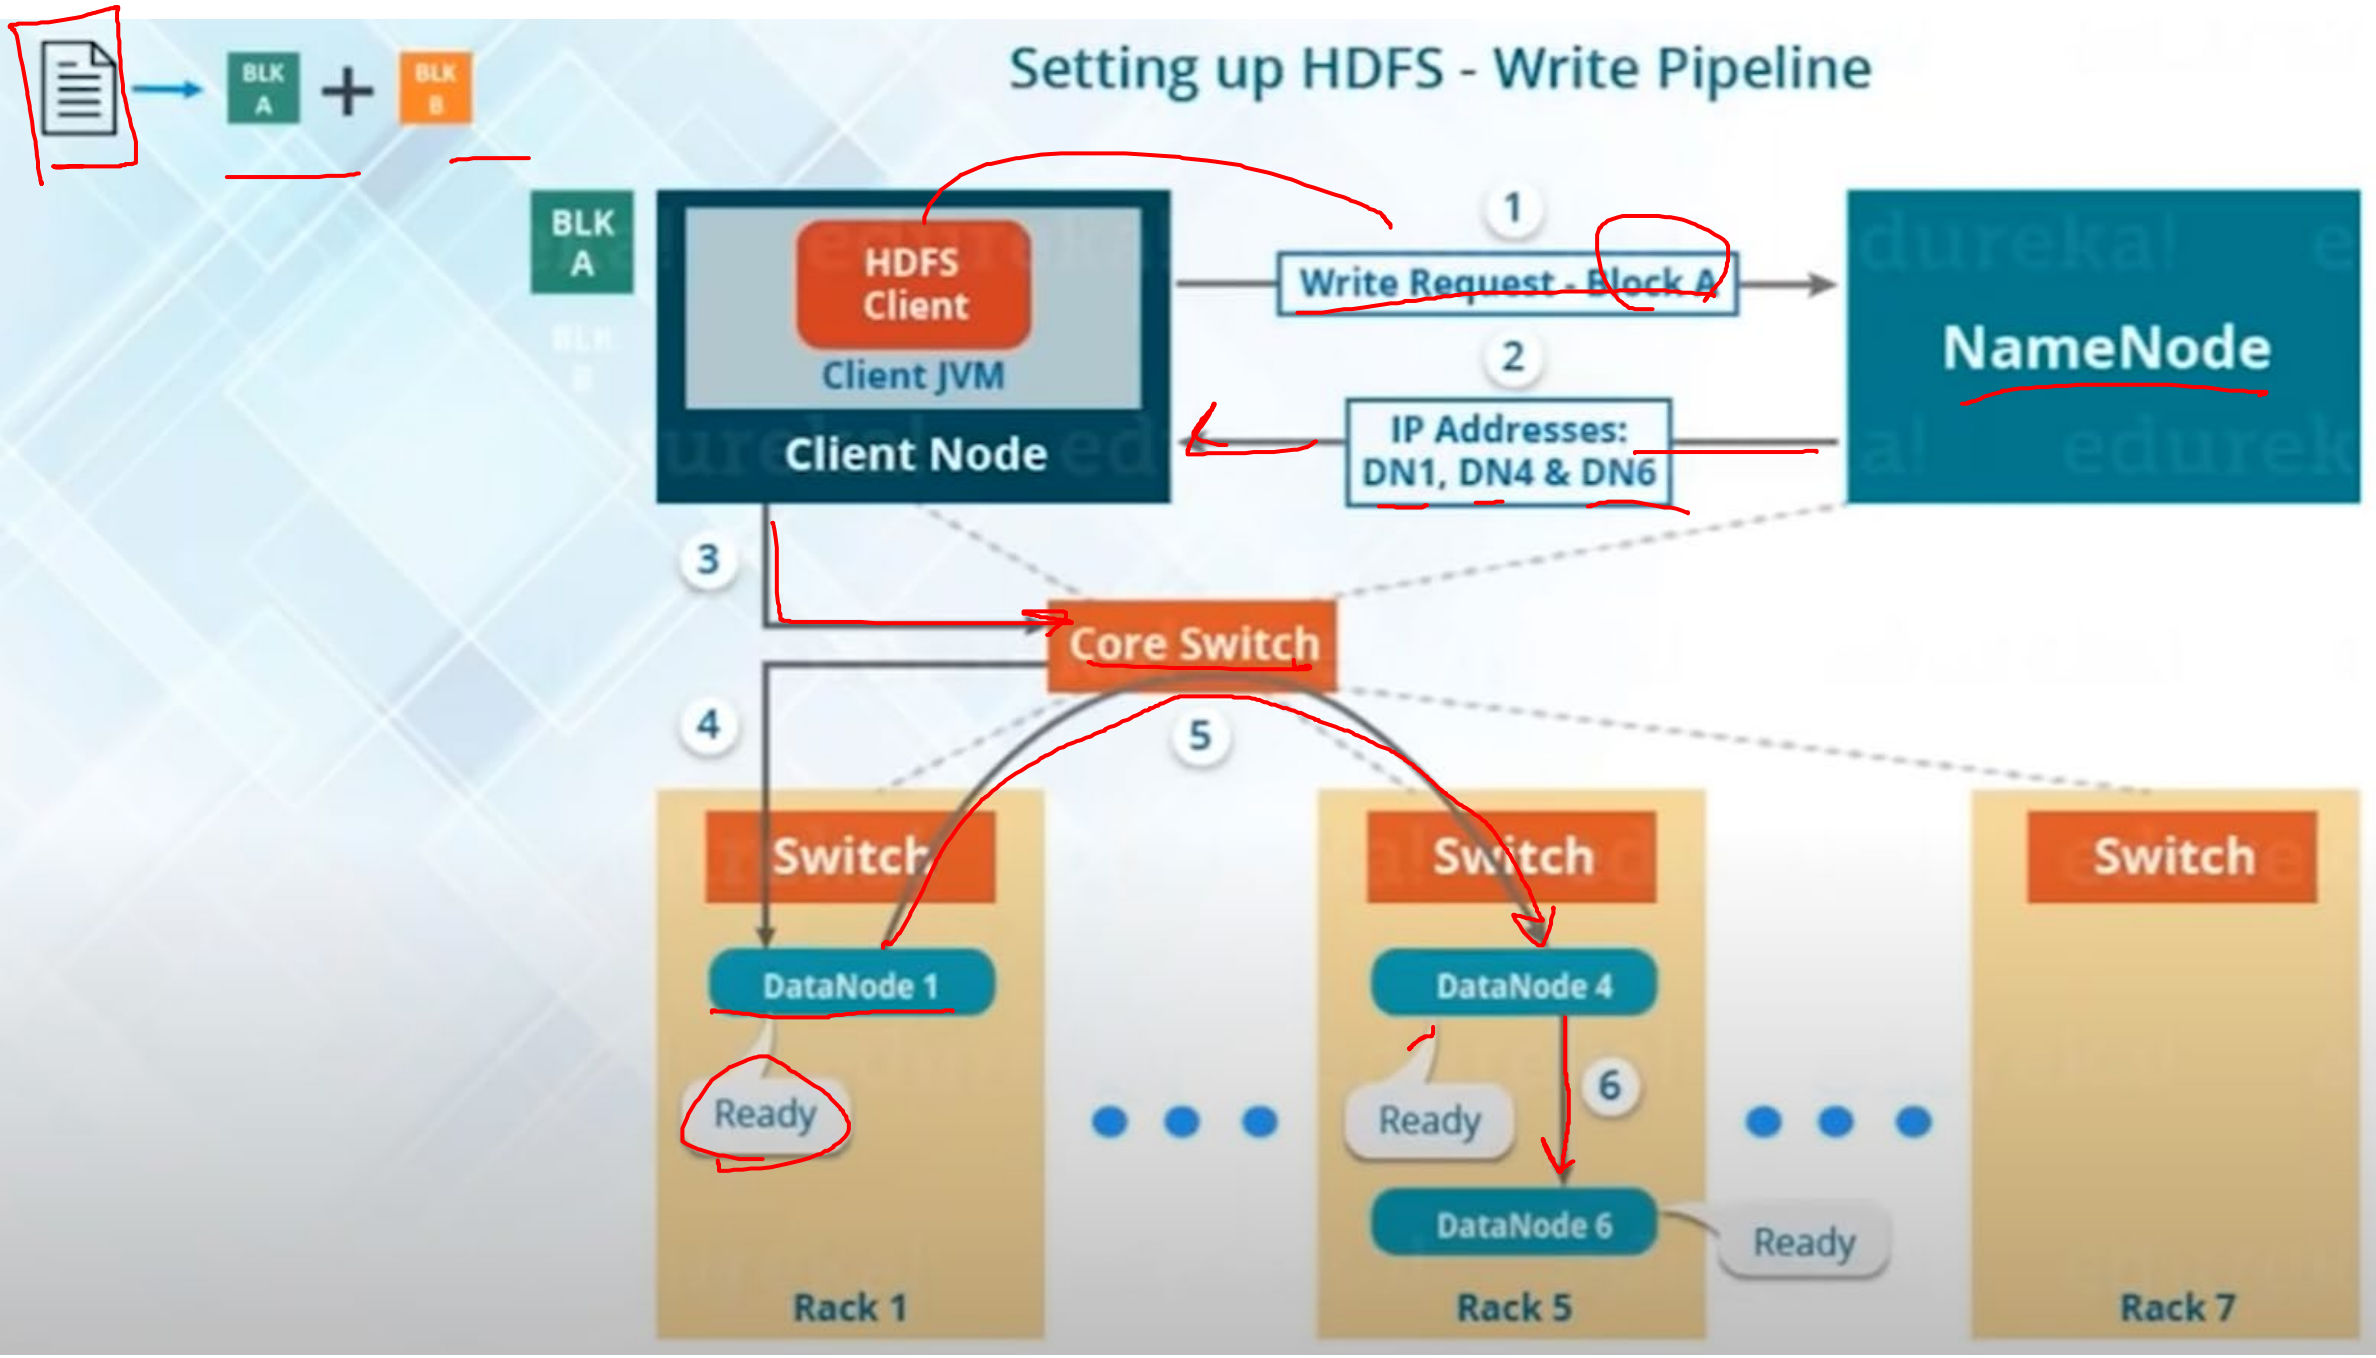

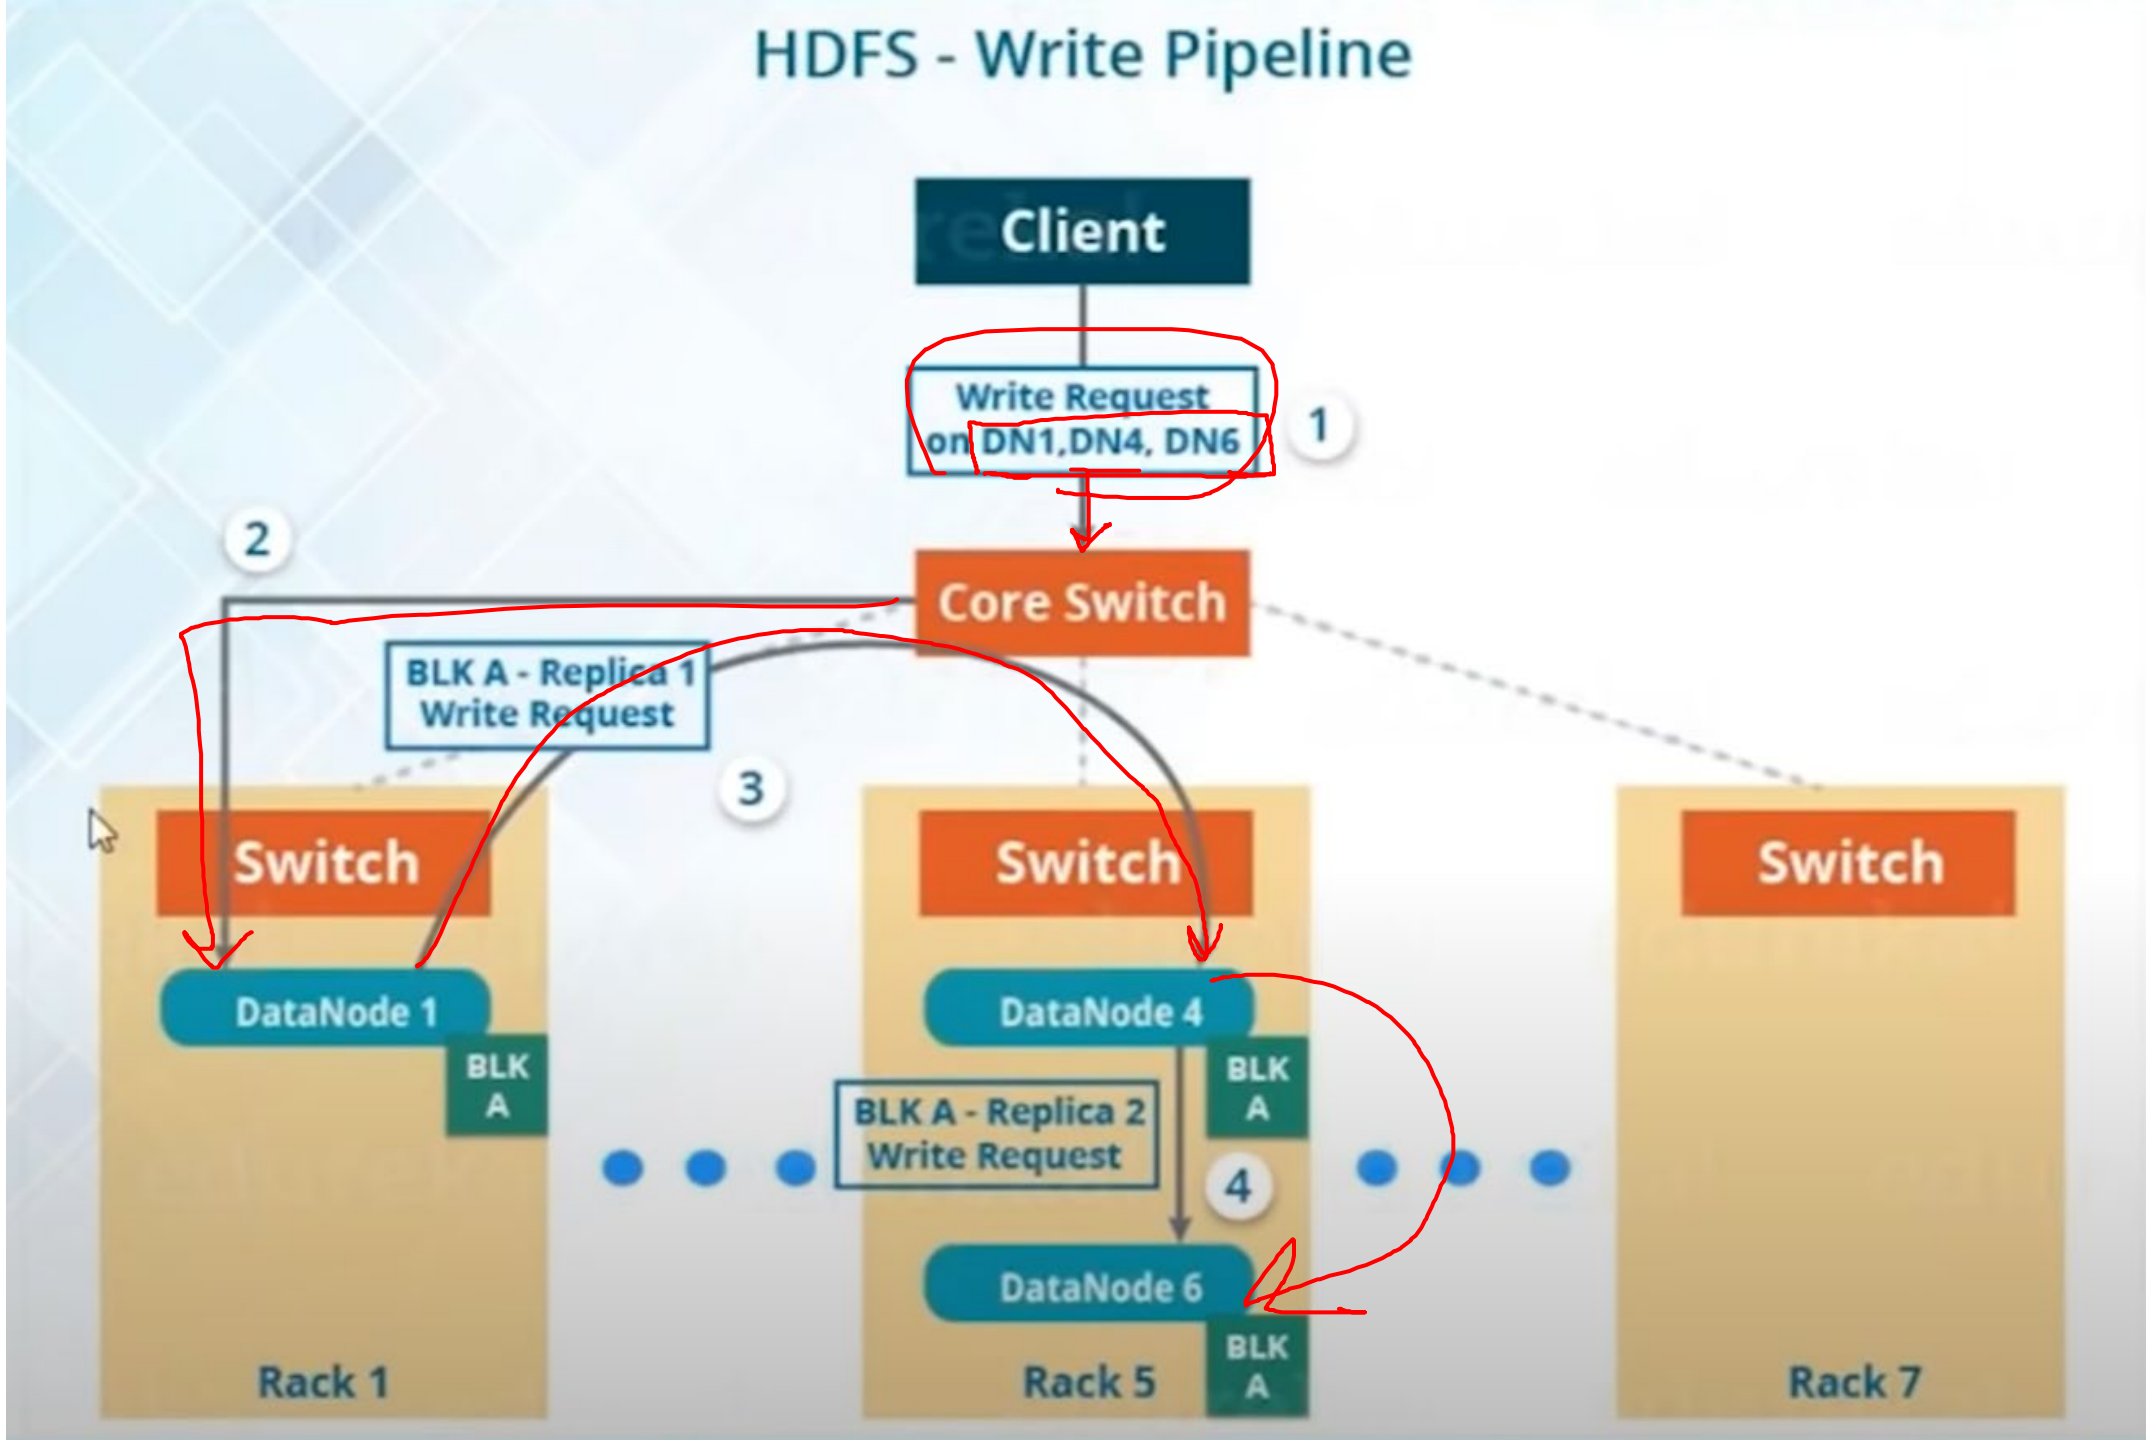
HDFS - Write Pipeline
Client
Write Request on DN1,DN4, DN6
1
2
Core Switch
BLK A - Replica 1 Write Request
3
Switch
Switch
Switch
DataNode 1
BLK A
DataNode 4
BLK A
BLK A - Replica 2 Write Request
4
DataNode 6
BLK A
Rack 1
Rack 5
Rack 7

Acknowledgement in HDFS - Write

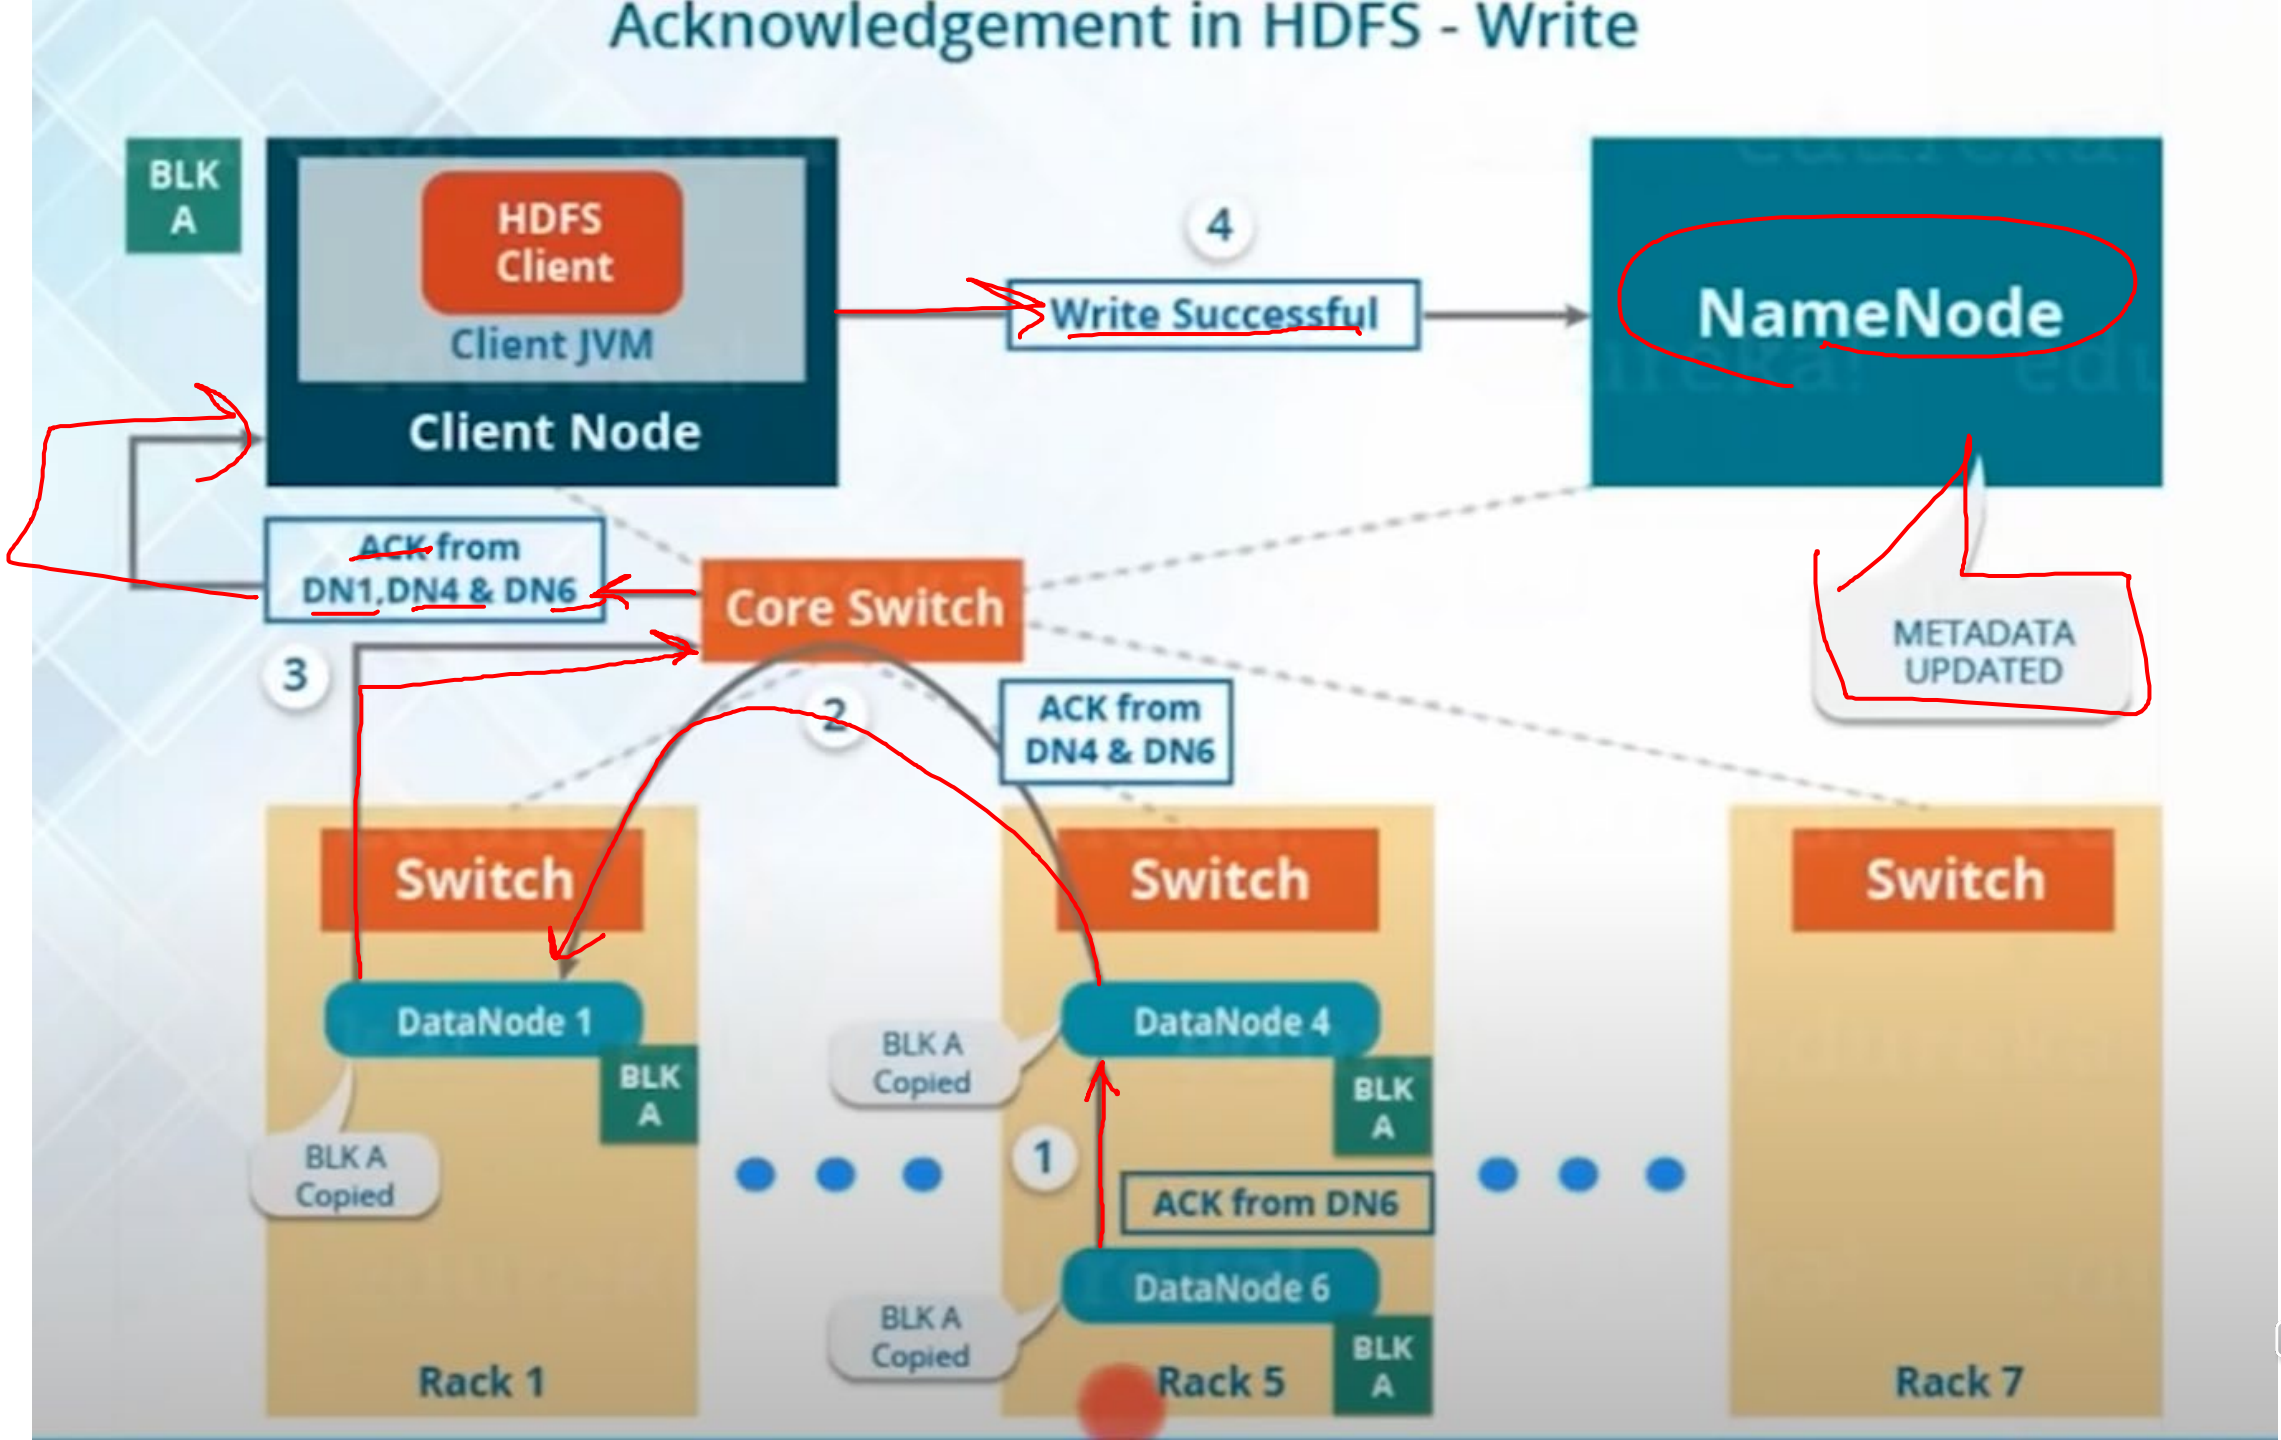

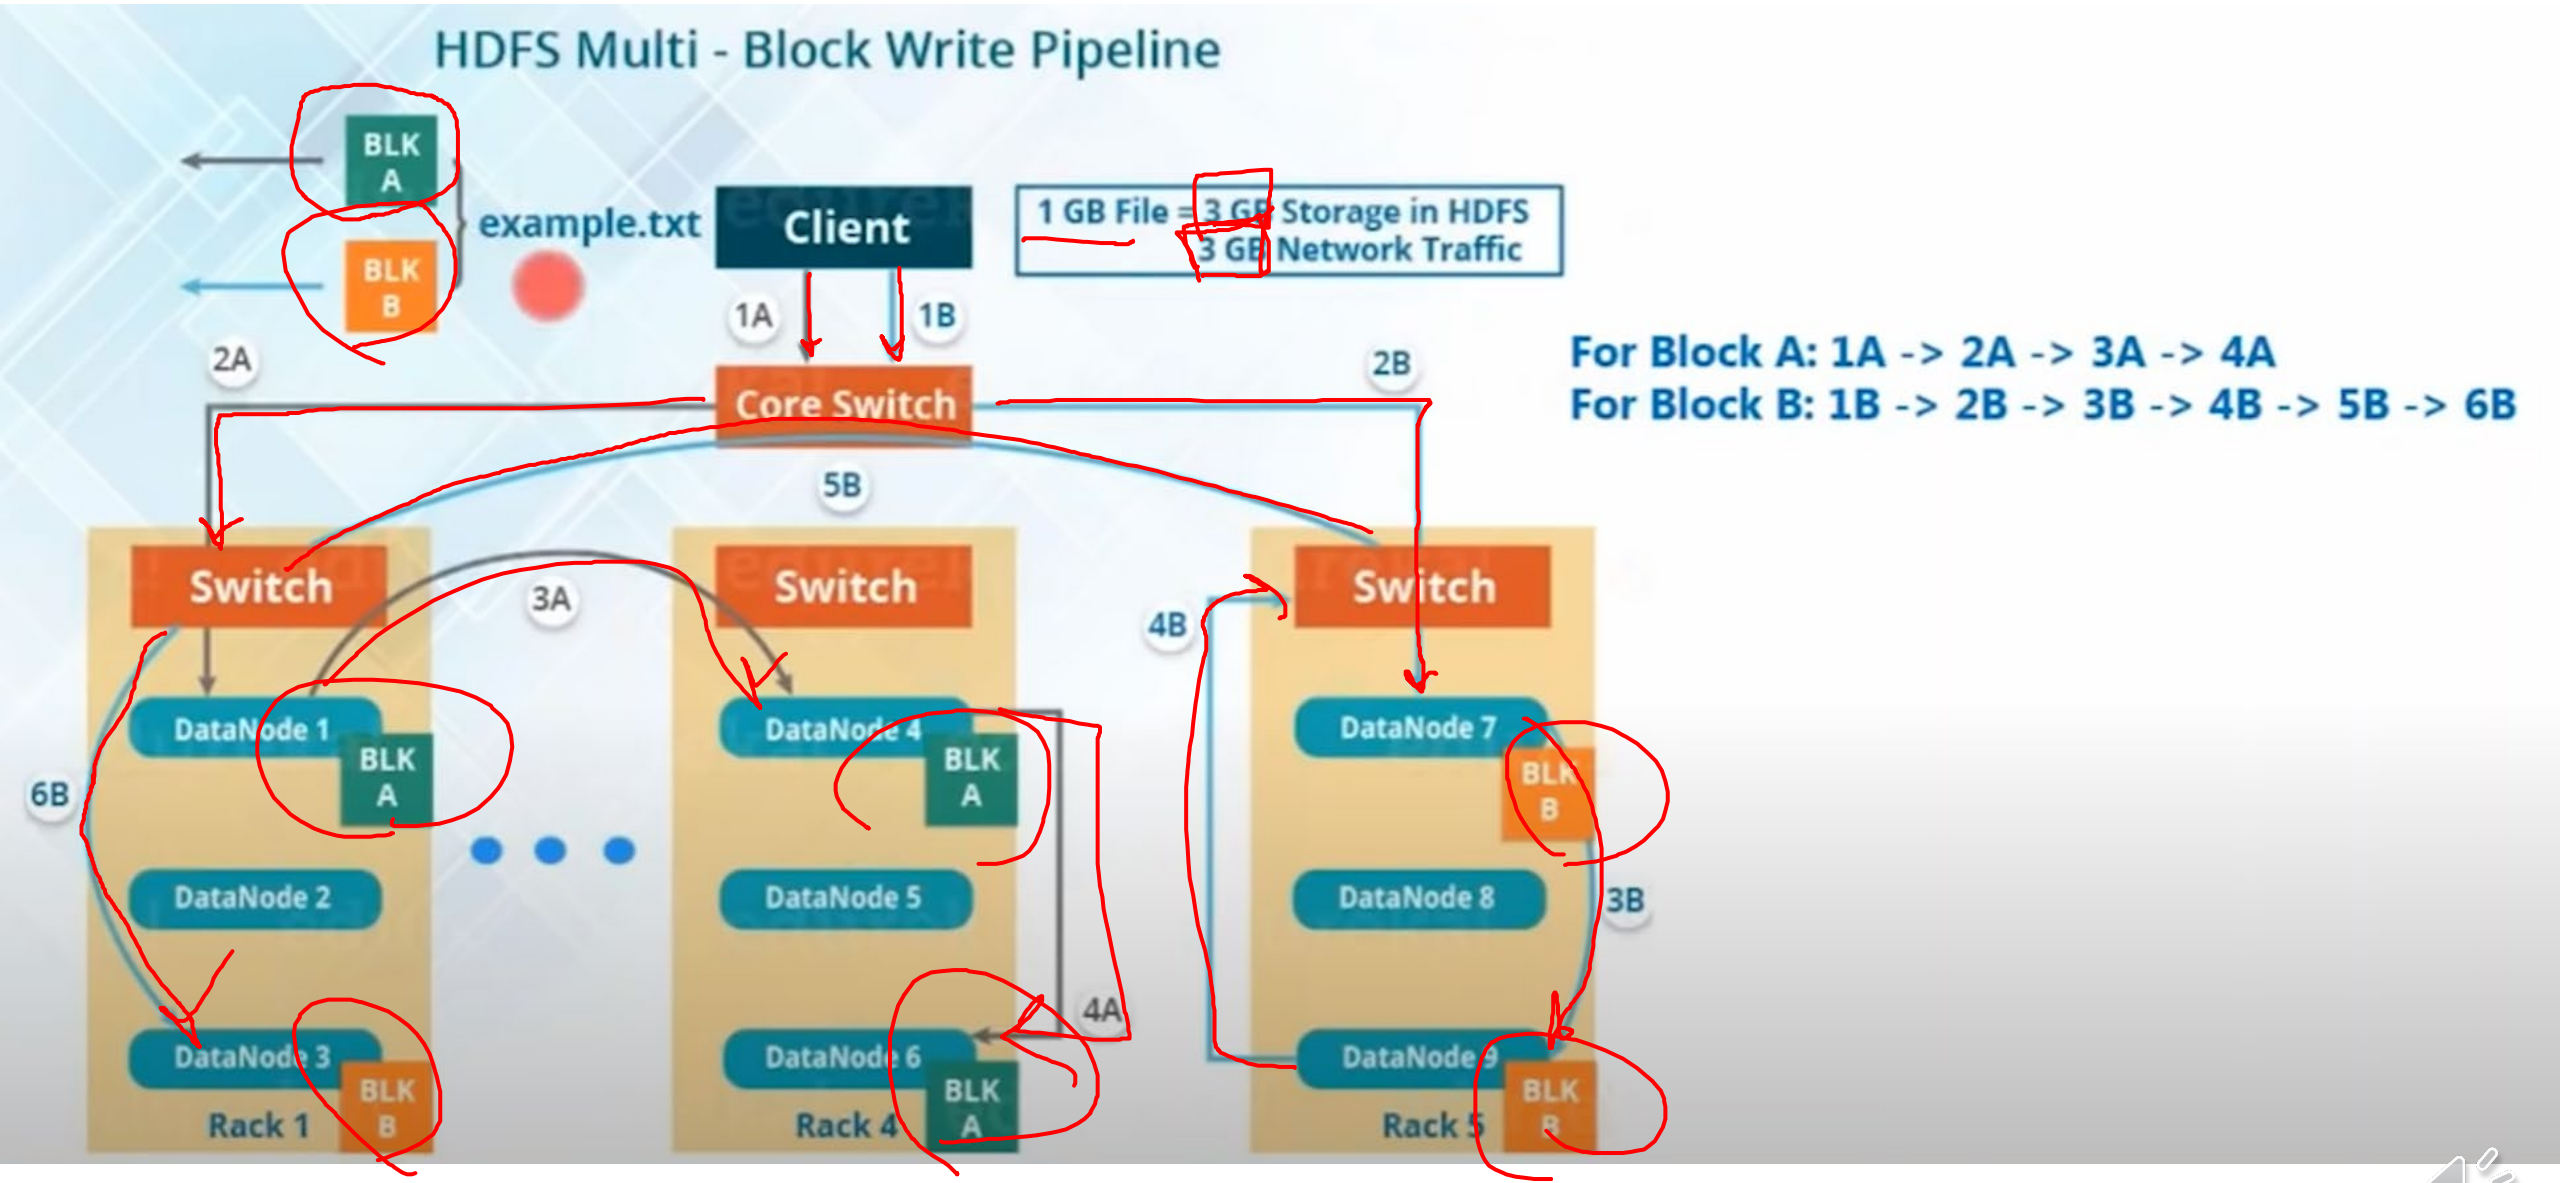
HDFS Multi - Block Write Pipeline
BLK A
BLK B
example.txt
Client
1 GB File = 3 GB Storage in HDFS
3 GB Network Traffic
1A
1B
2A
2B
Core Switch
5B
Switch
Switch
Switch
3A
4B
6B
DataNode 1
BLK A
DataNode 2
DataNode 3
BLK B
Rack 1
DataNode 4
BLK A
DataNode 5
DataNode 6
BLK A
4A
Rack 4
DataNode 7
BLK B
DataNode 8
3B
DataNode 9
BLK B
Rack 5
For Block A: 1A -> 2A -> 3A -> 4A
For Block B: 1B -> 2B -> 3B -> 4B -> 5B -> 6B

HDFS - Read Architecture

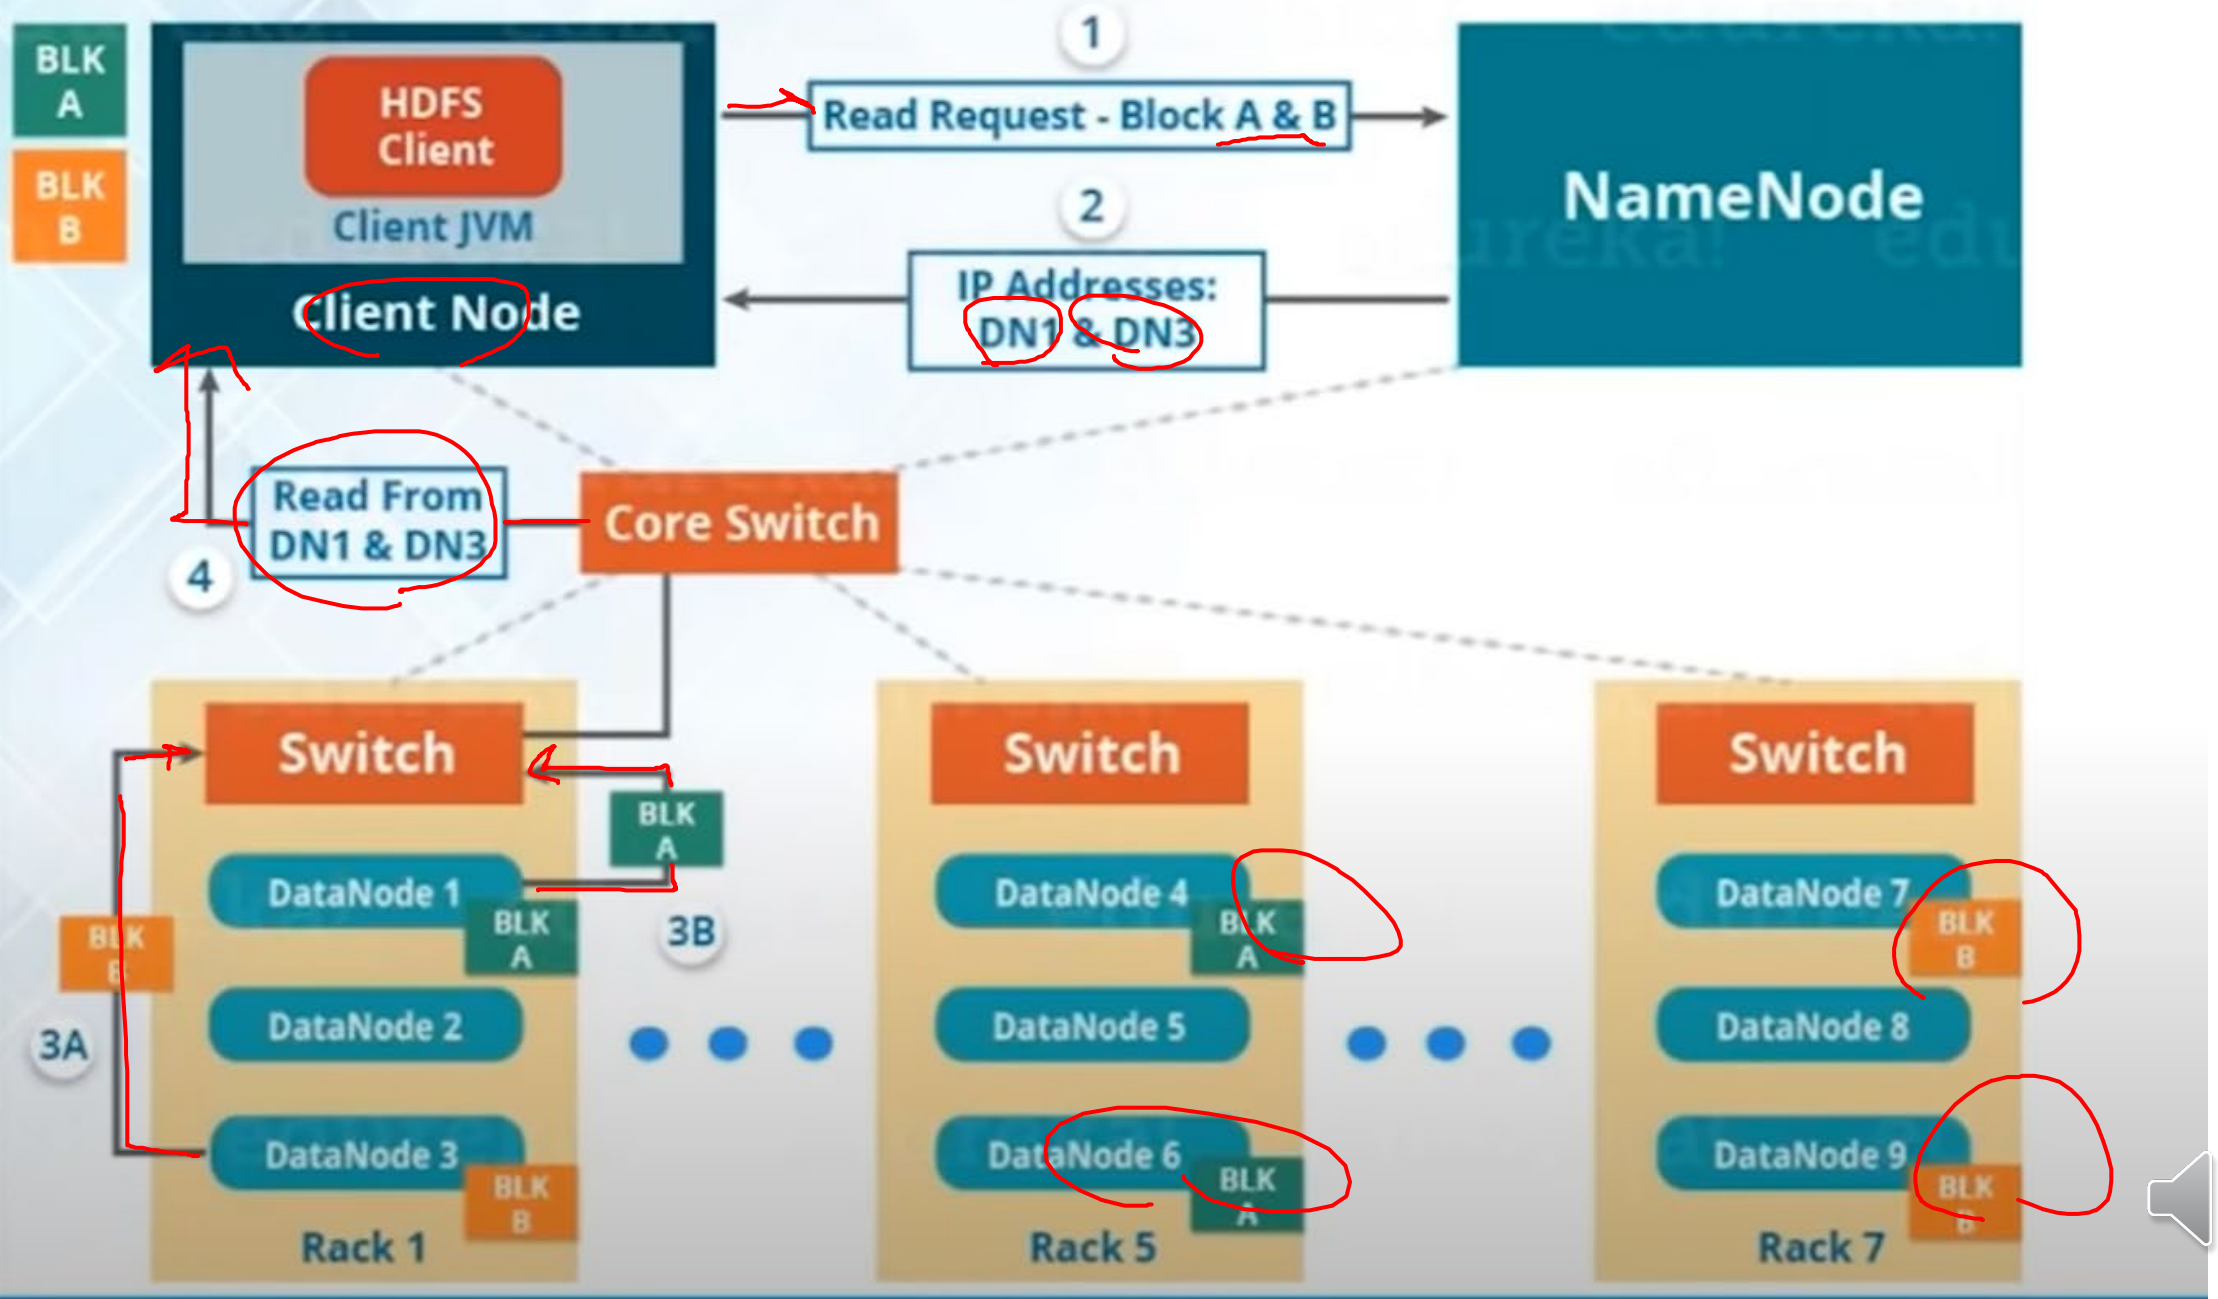

HADOOP CORE COMPONENTS

HDFS

Storage:
Distributed File System

MapReduce

Processing:
Allows parallel & distributed processing

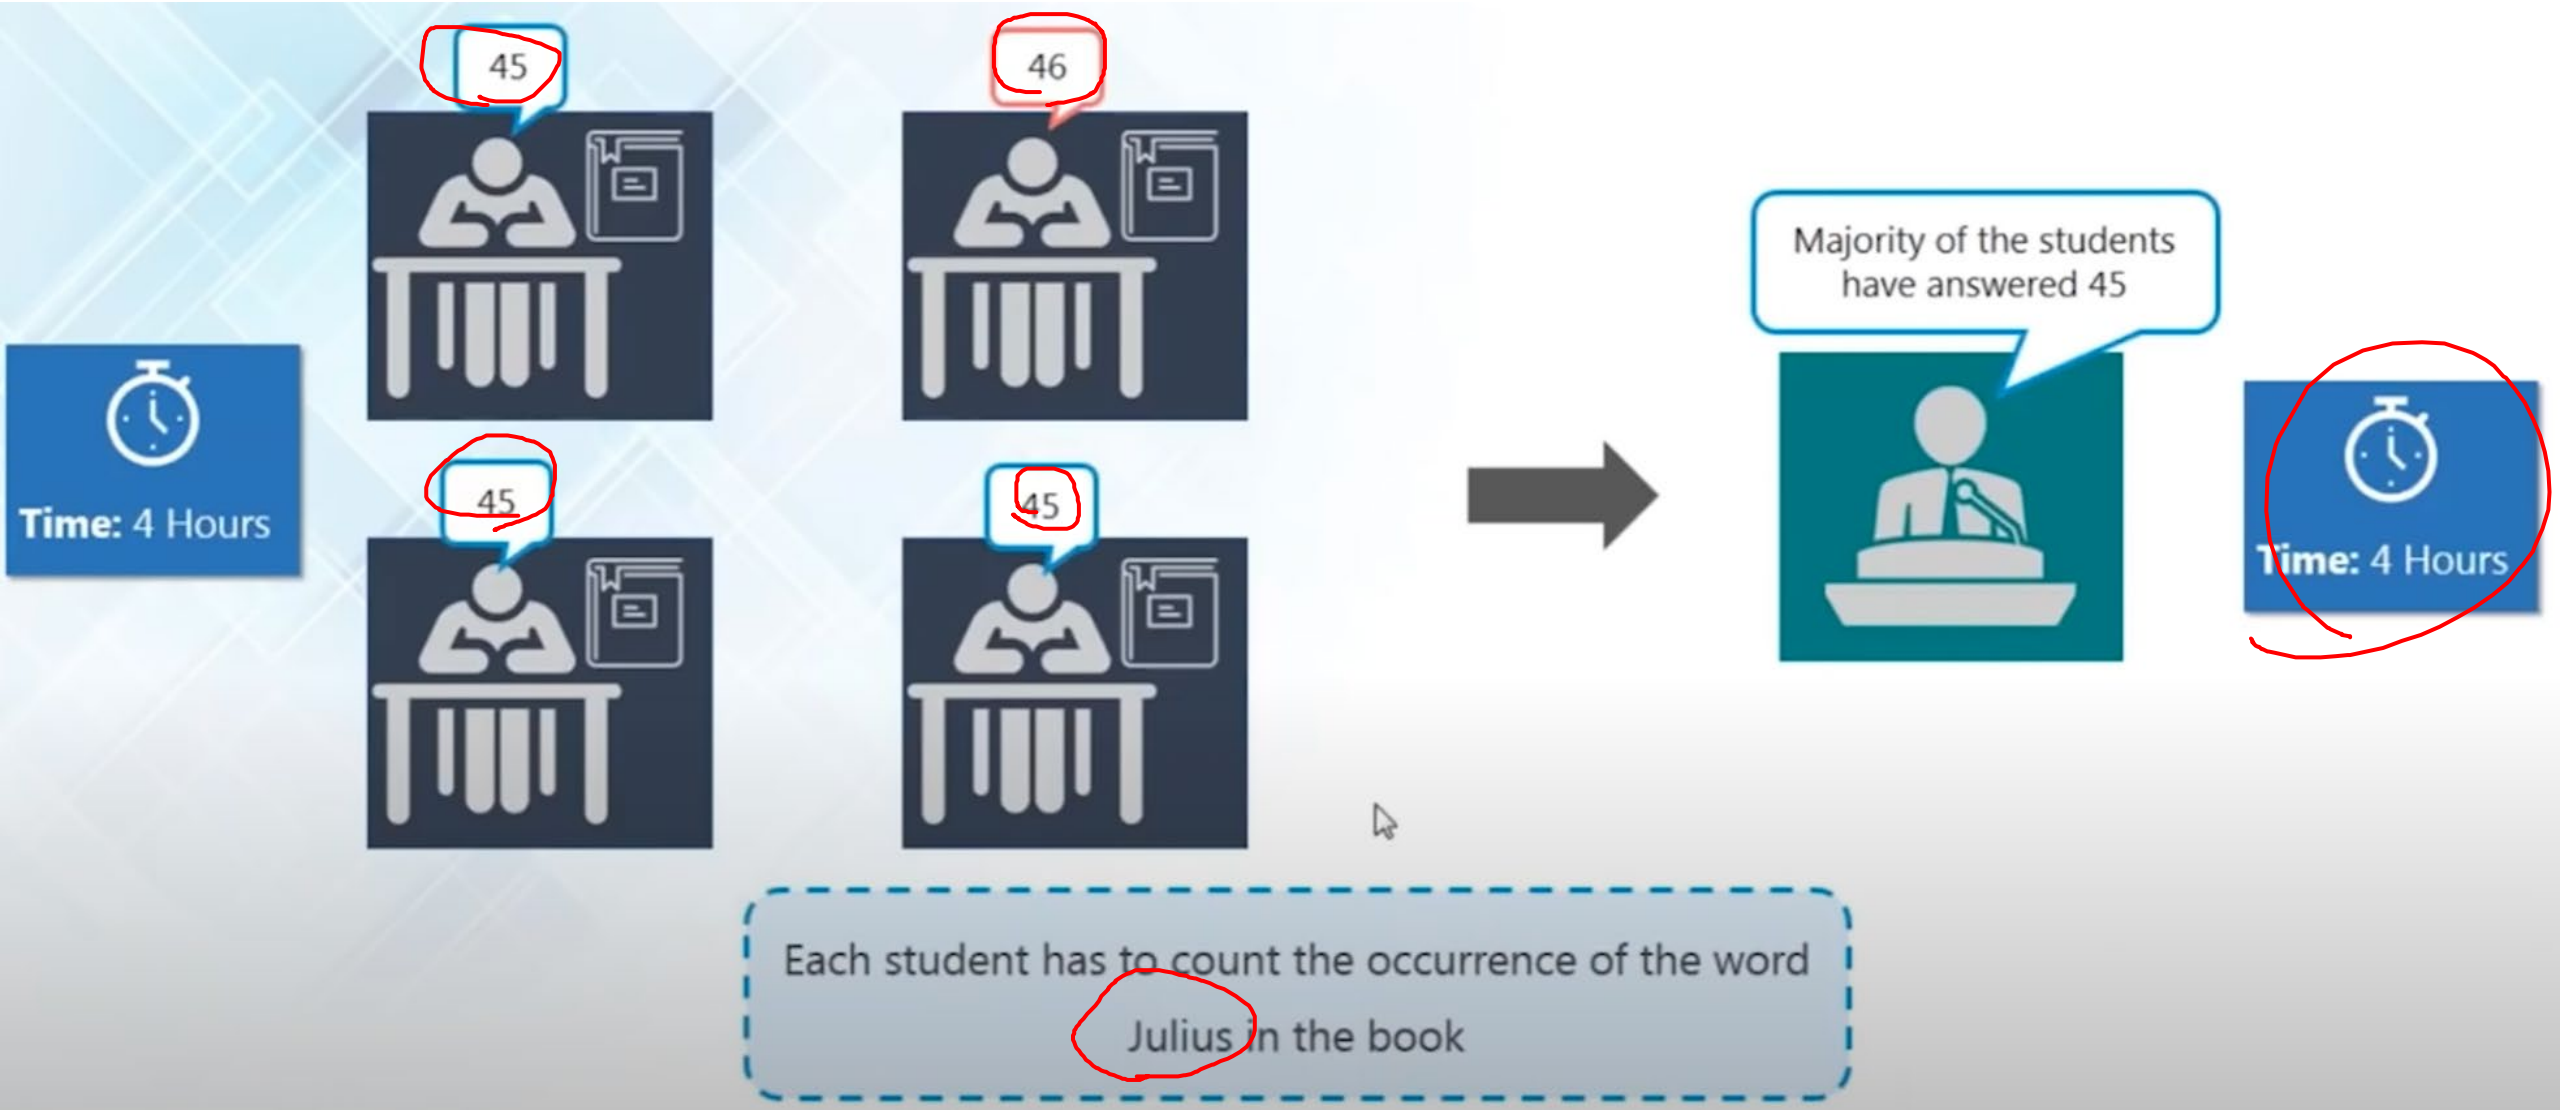
Time: 4 Hours
45
46
45
45
Majority of the students have answered 45
Time: 4 Hours
Each student has to count the occurrence of the word Julius in the book

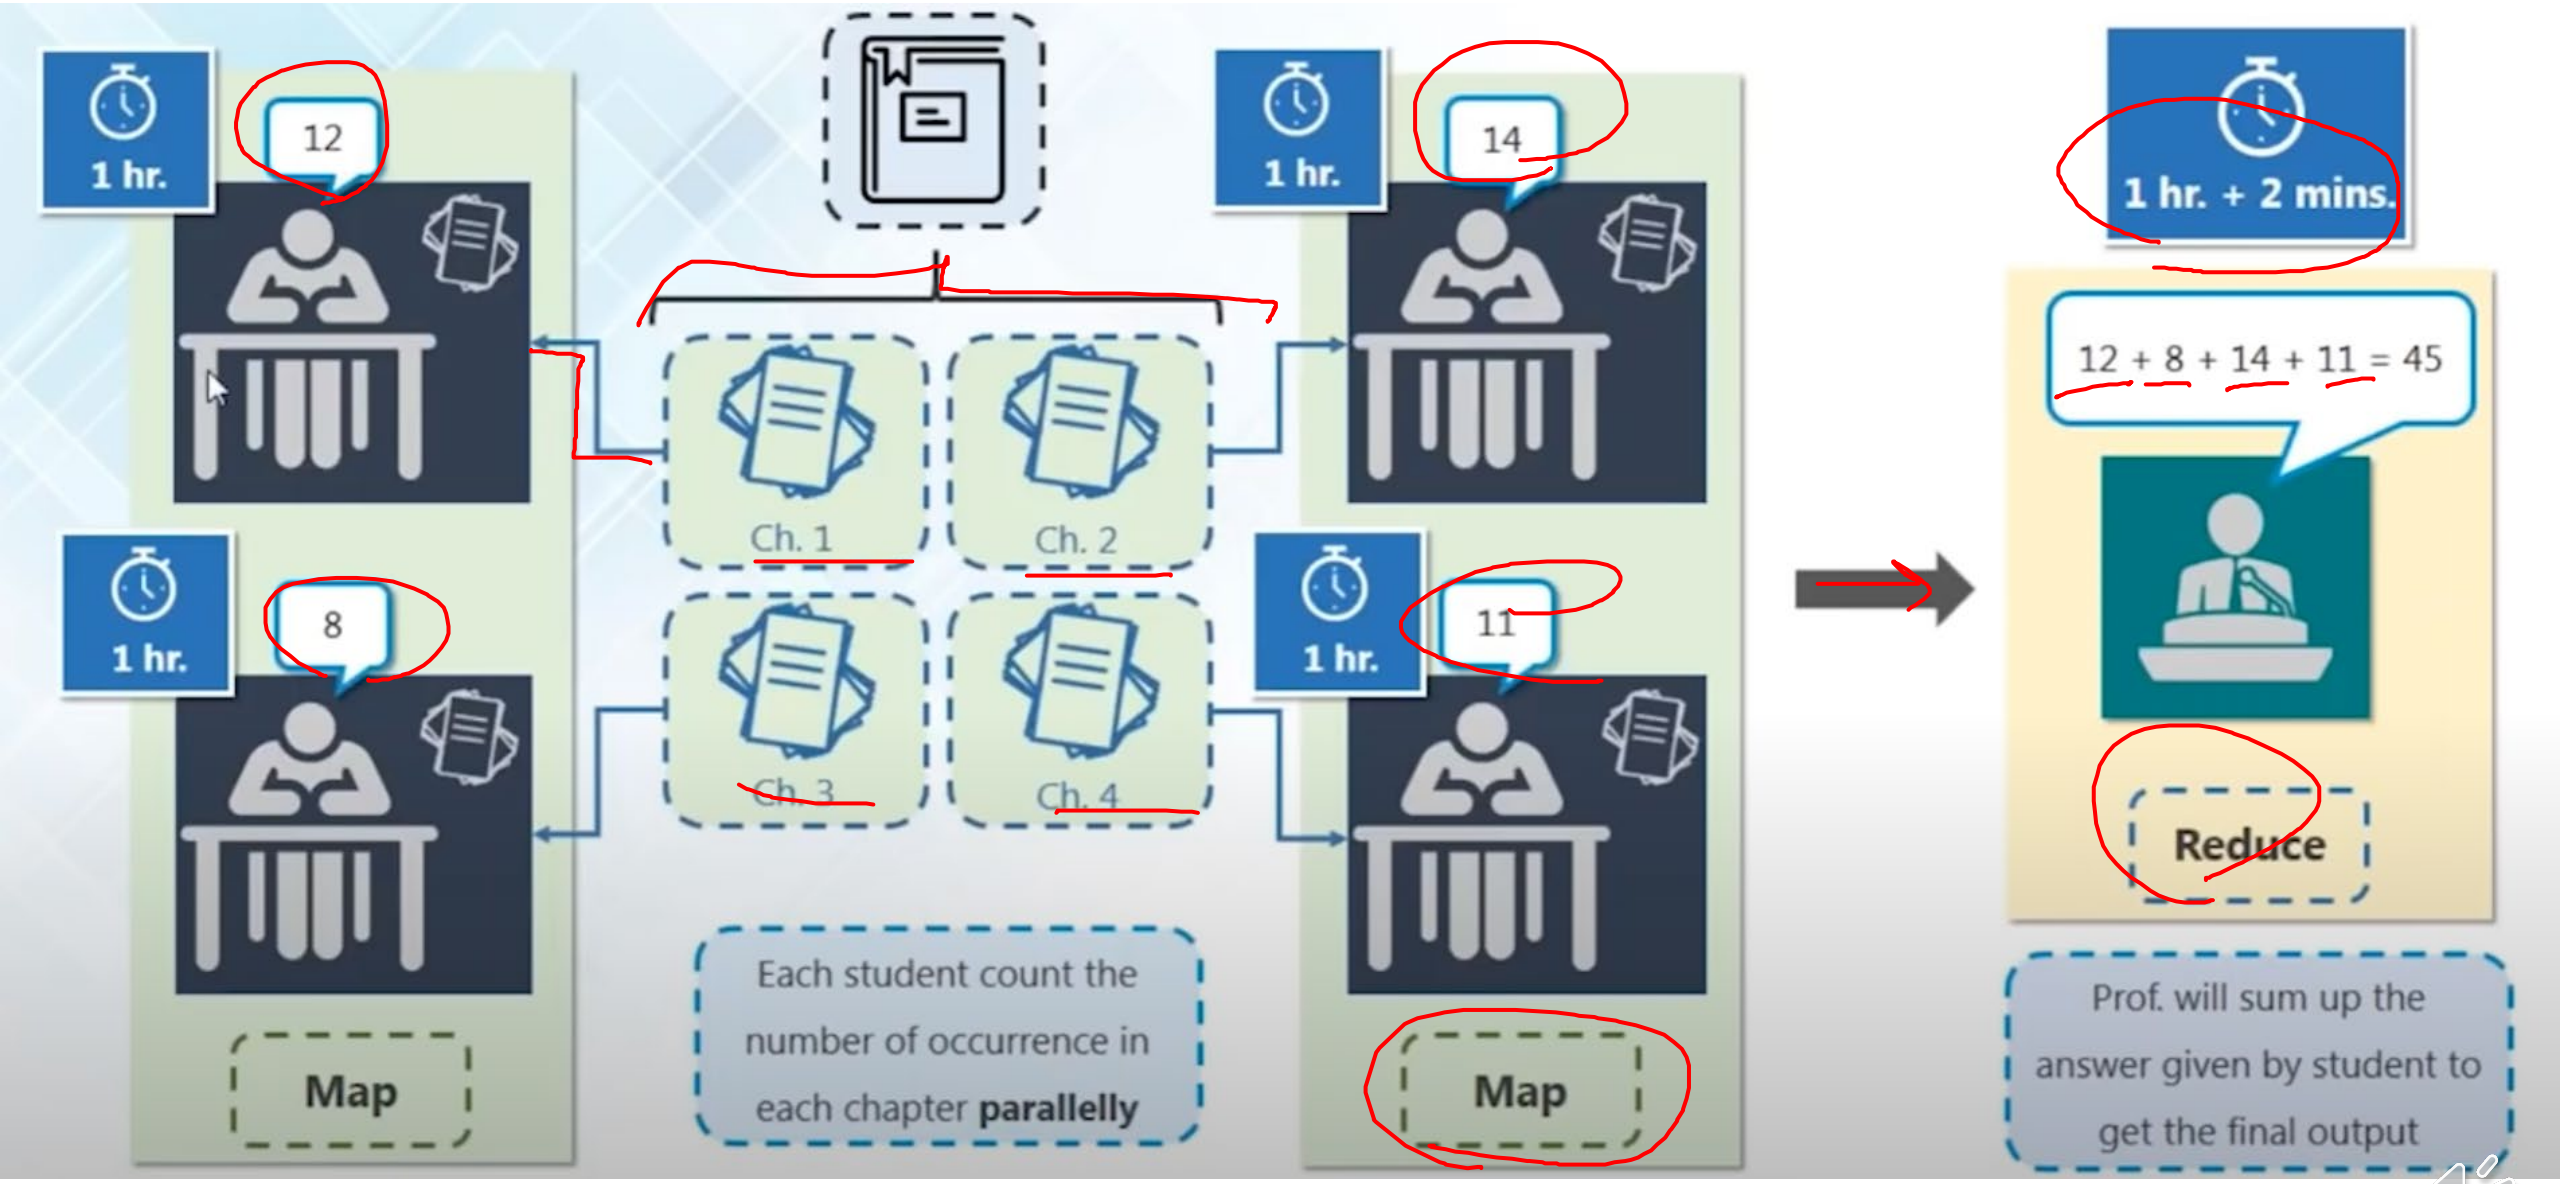
1 hr.
12
1 hr.
14
Ch. 1
Ch. 2
Ch. 3
Ch. 4
1 hr.
8
1 hr.
11
Map
Each student count the number of occurrence in each chapter **parallelly**
Map
1 hr. + 2 mins.
12 + 8 + 14 + 11 = 45
Reduce
Prof. will sum up the answer given by student to get the final output

MapReduce is a programming framework that allows us to perform distributed and parallel processing on large data sets in a distributed environment

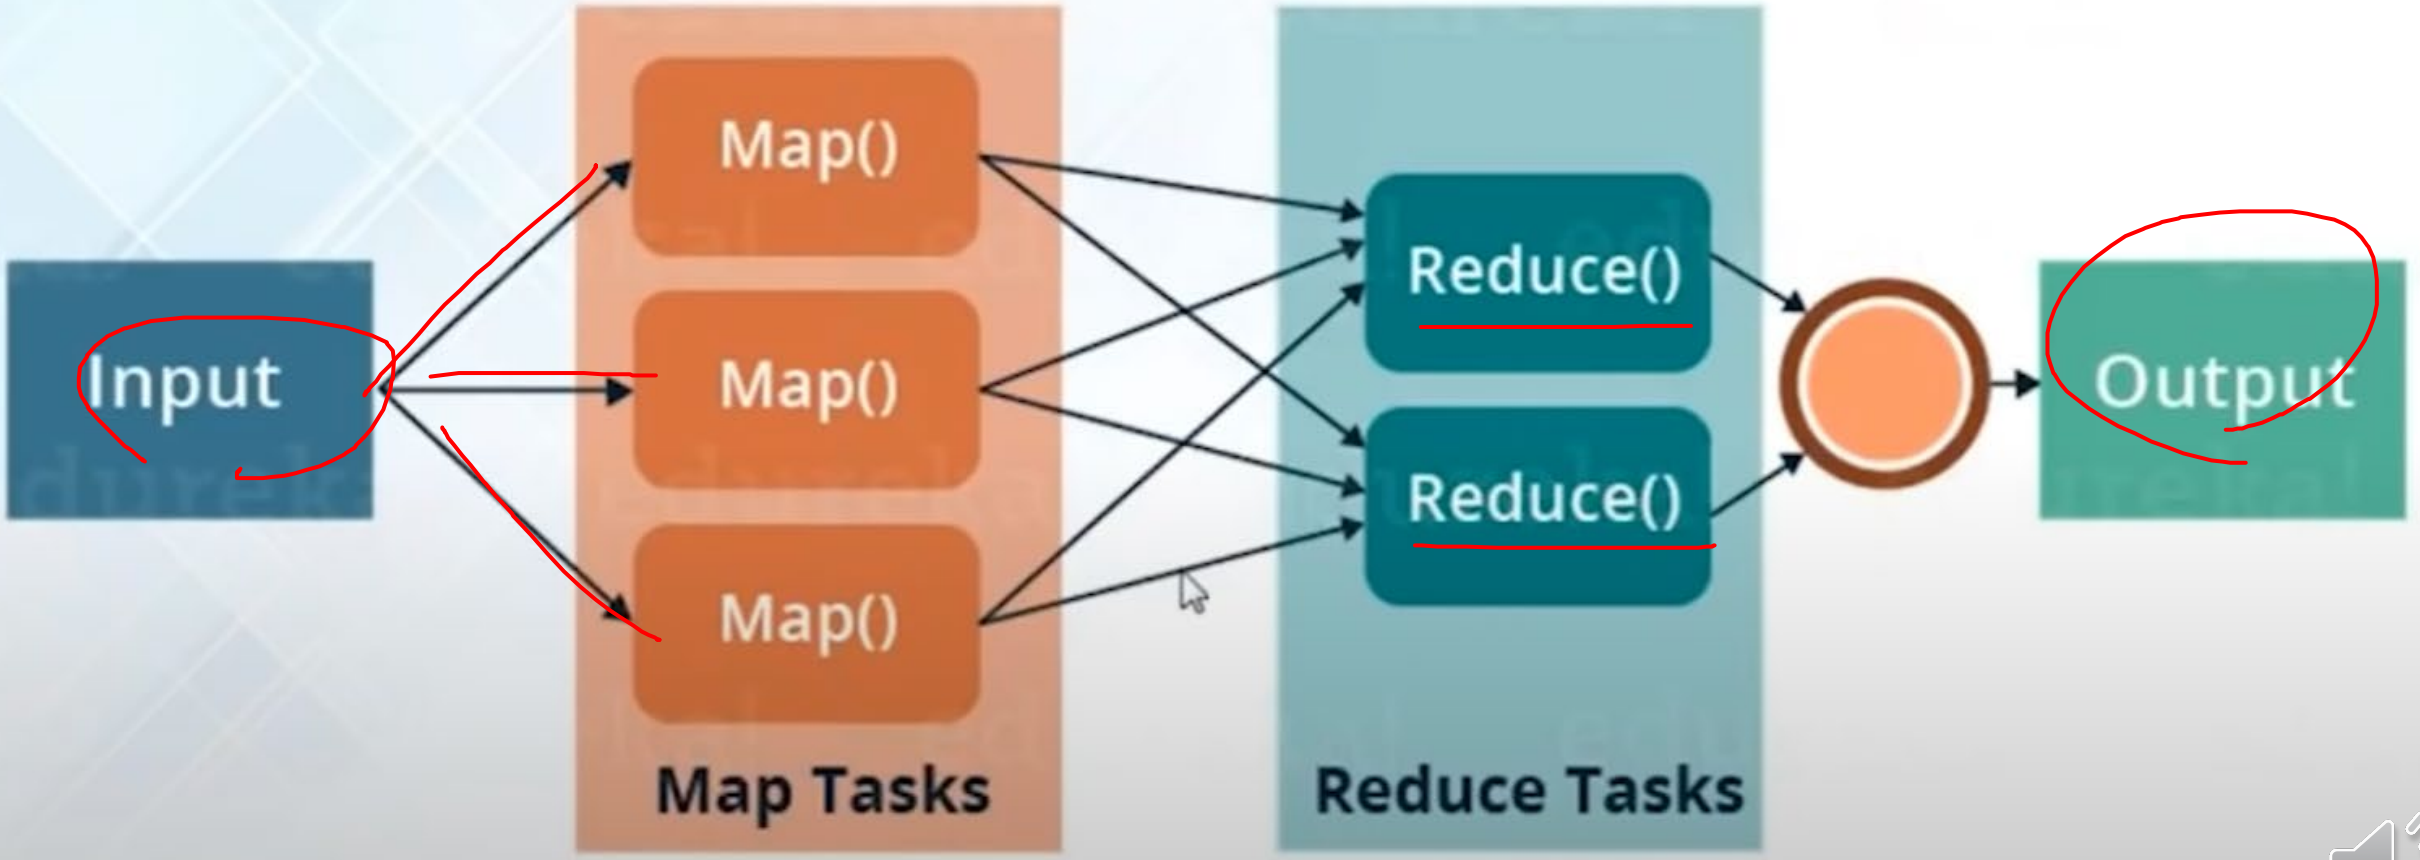

MapReduce
Word Count Program

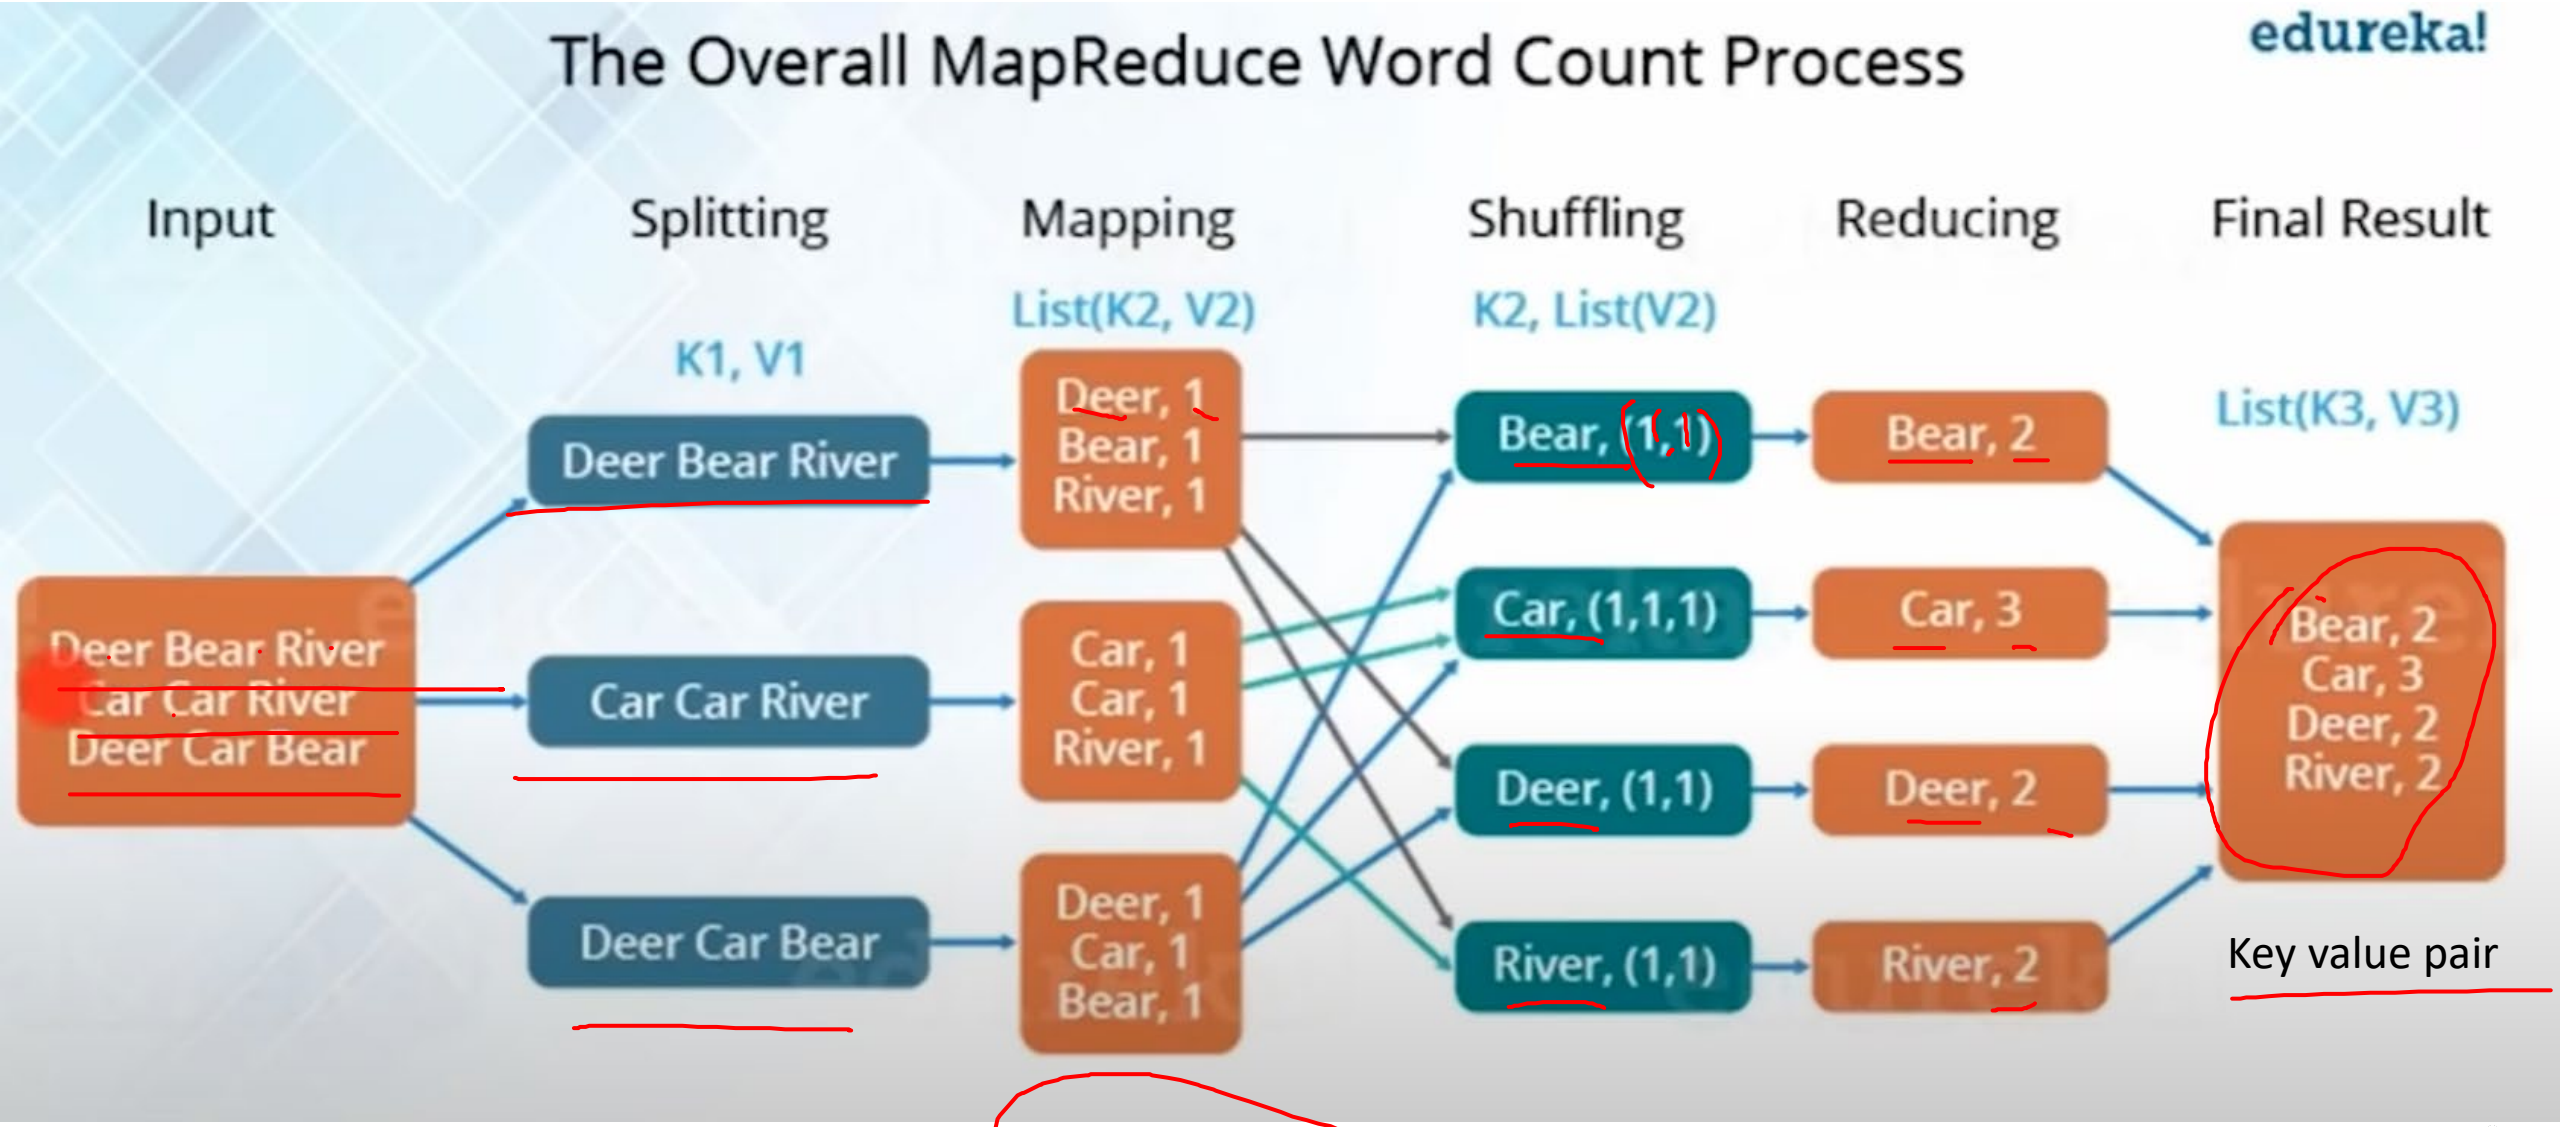
The Overall MapReduce Word Count Process
edureka!
Input
Splitting
Mapping
Shuffling
Reducing
Final Result
K1, V1
List(K2, V2)
K2, List(V2)
List(K3, V3)
Deer Bear River
Car Car River
Deer Car Bear
Deer Bear River
Car Car River
Deer Car Bear
Deer, 1
Bear, 1
River, 1
Car, 1
Car, 1
River, 1
Deer, 1
Car, 1
Bear, 1
Bear, (1,1)
Car, (1,1,1)
Deer, (1,1)
River, (1,1)
Bear, 2
Car, 3
Deer, 2
River, 2
Bear, 2
Car, 3
Deer, 2
River, 2
Key value pair
Key value pair

Three Major Parts of MapReduce Program:

1

Mapper Code:

You write the mapper logic over here i.e. how map task will process the data to produce the key-value pair to be aggregated

2

Reducer Code:

You write reducer logic here which combines the intermediate key-value pair generated by Mapper to give the final aggregated output

3

Driver Code

You specify all the job configurations over here like job name, Input path, output path, etc.